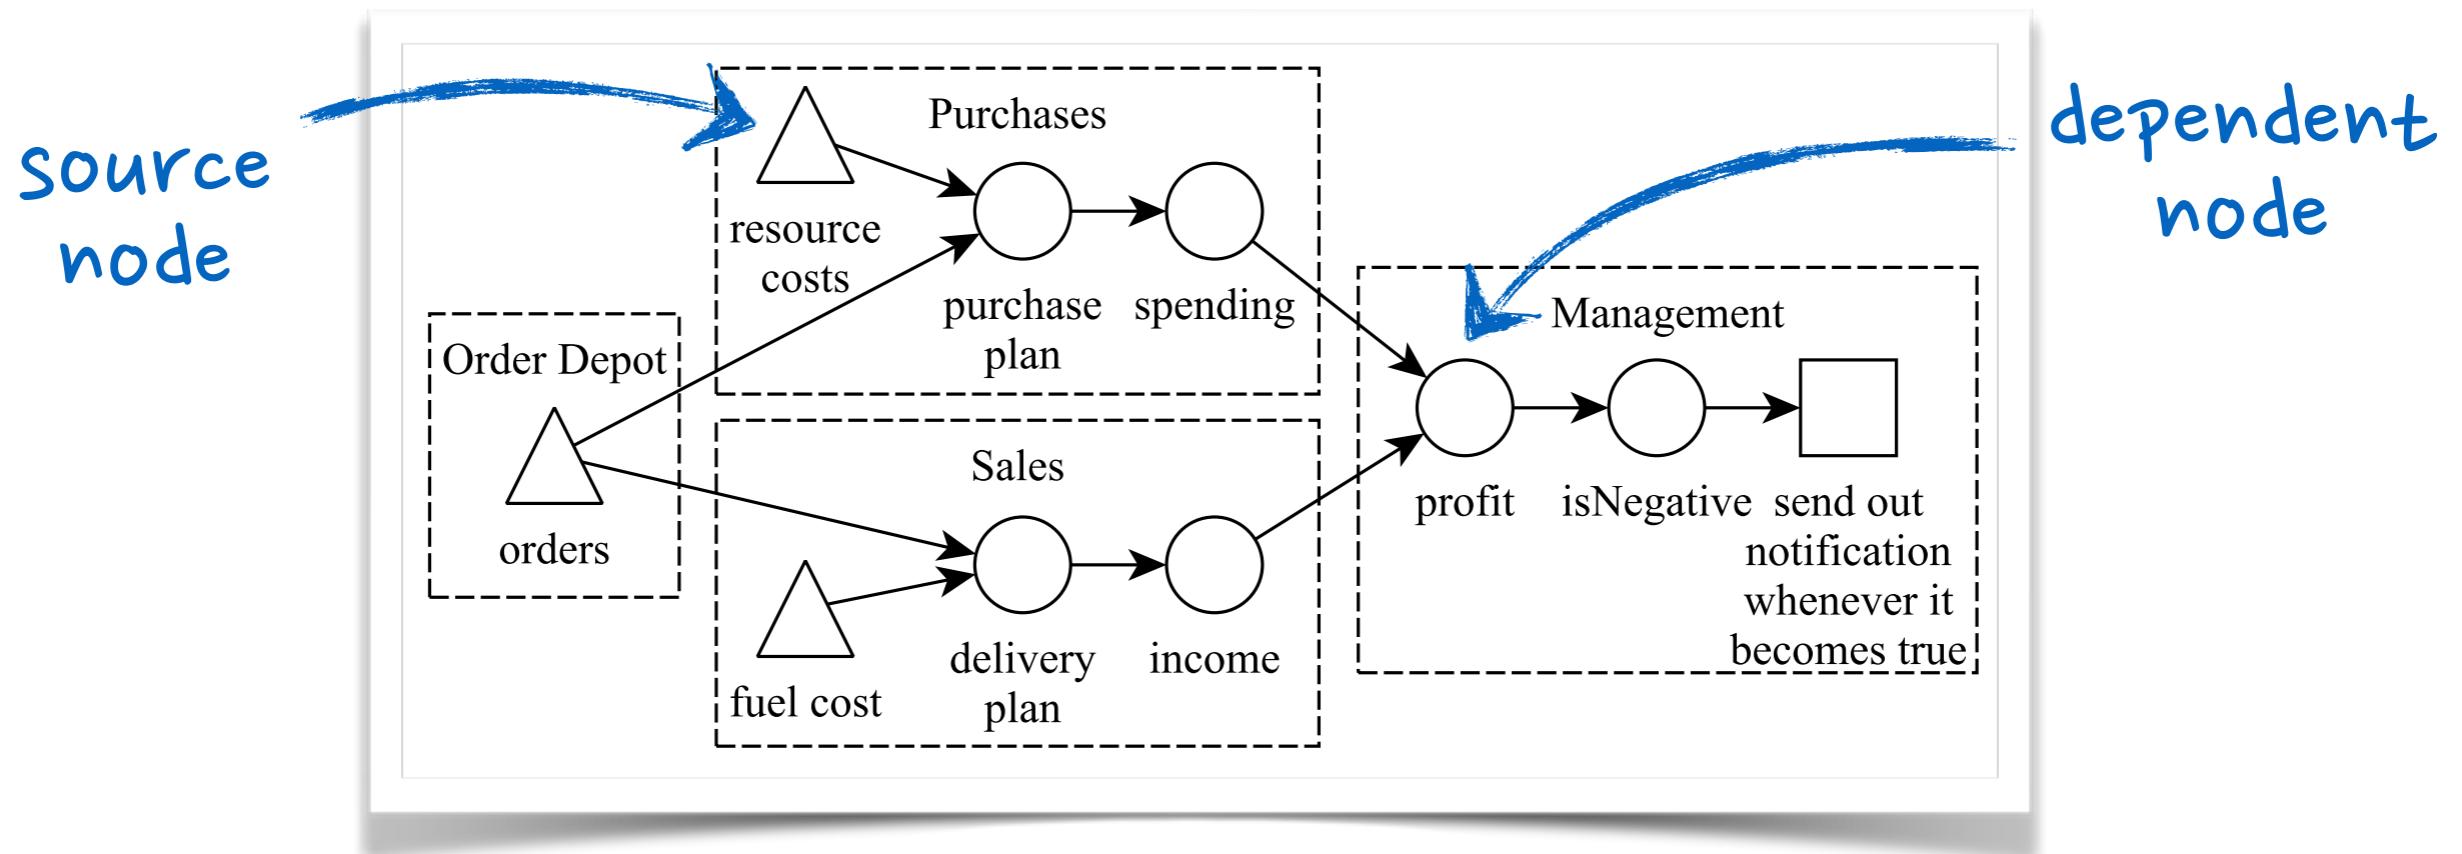


Distributed REScala: an update algorithm for distributed Reactive Programming

(José Proença @ 1 April)

OOPSLA'14 - Joscha Drechsler, Guido Salvaneschi,
Ragnar Mogk, and Mira Mezini
(TU-Darmstadt)

Context and Motivation



Lack of generic distributed algorithms

- only client-side, no glitch-freedom, or cannot be distributed

Distributed RP vs. observers + remote obj.

Contributions

Algorithms for distributed glitch-freedom:

Scala.React

Topological sorting + priority queue

Scala.Rx

Parallel propagator version

ELM

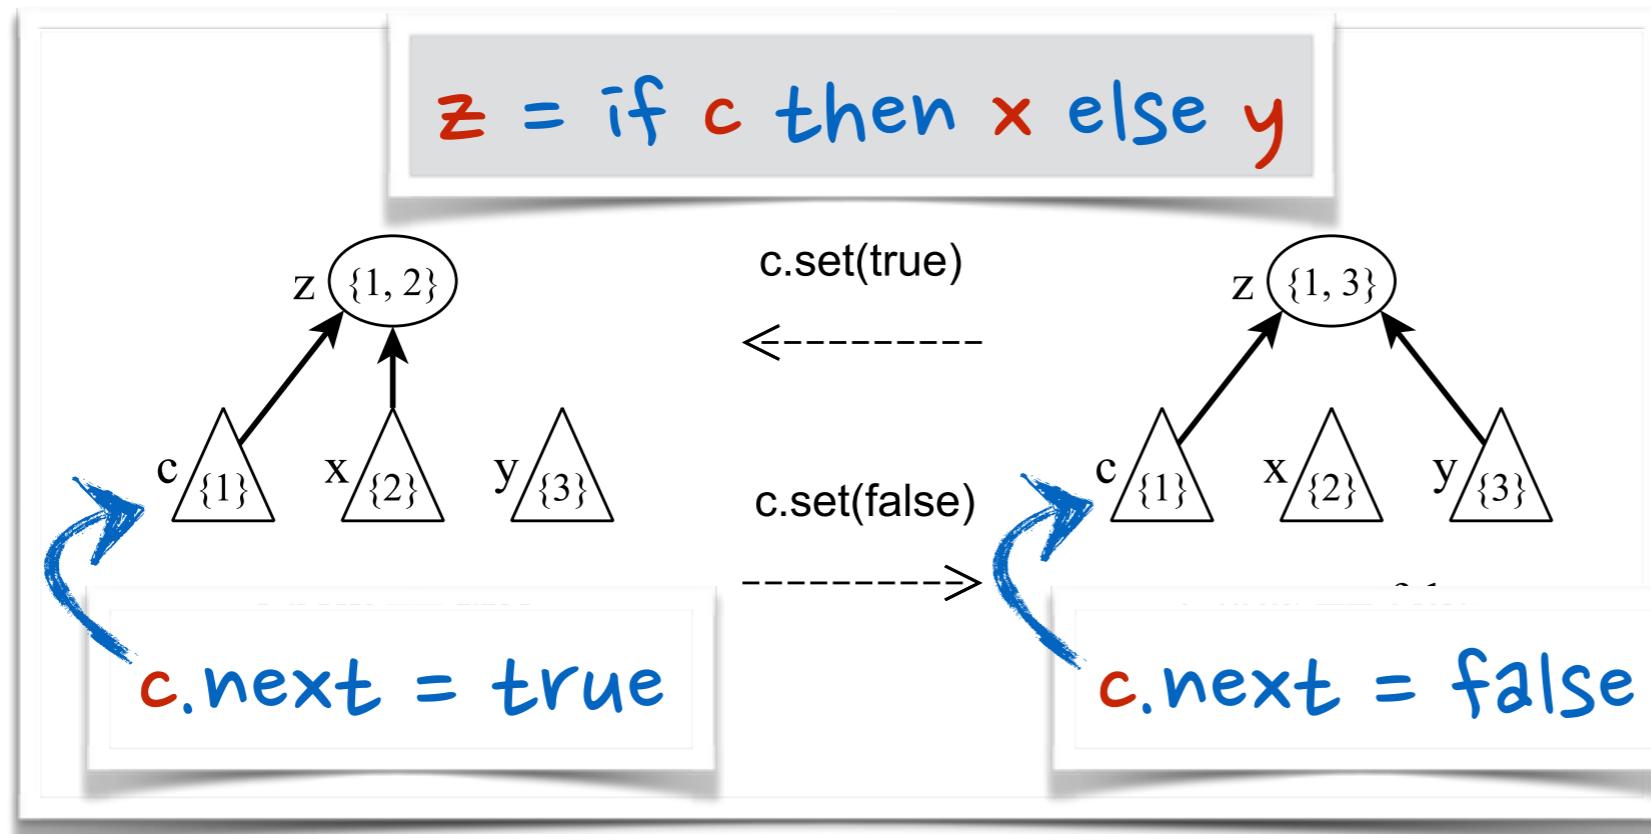
Decentralised flooding

SID-UP

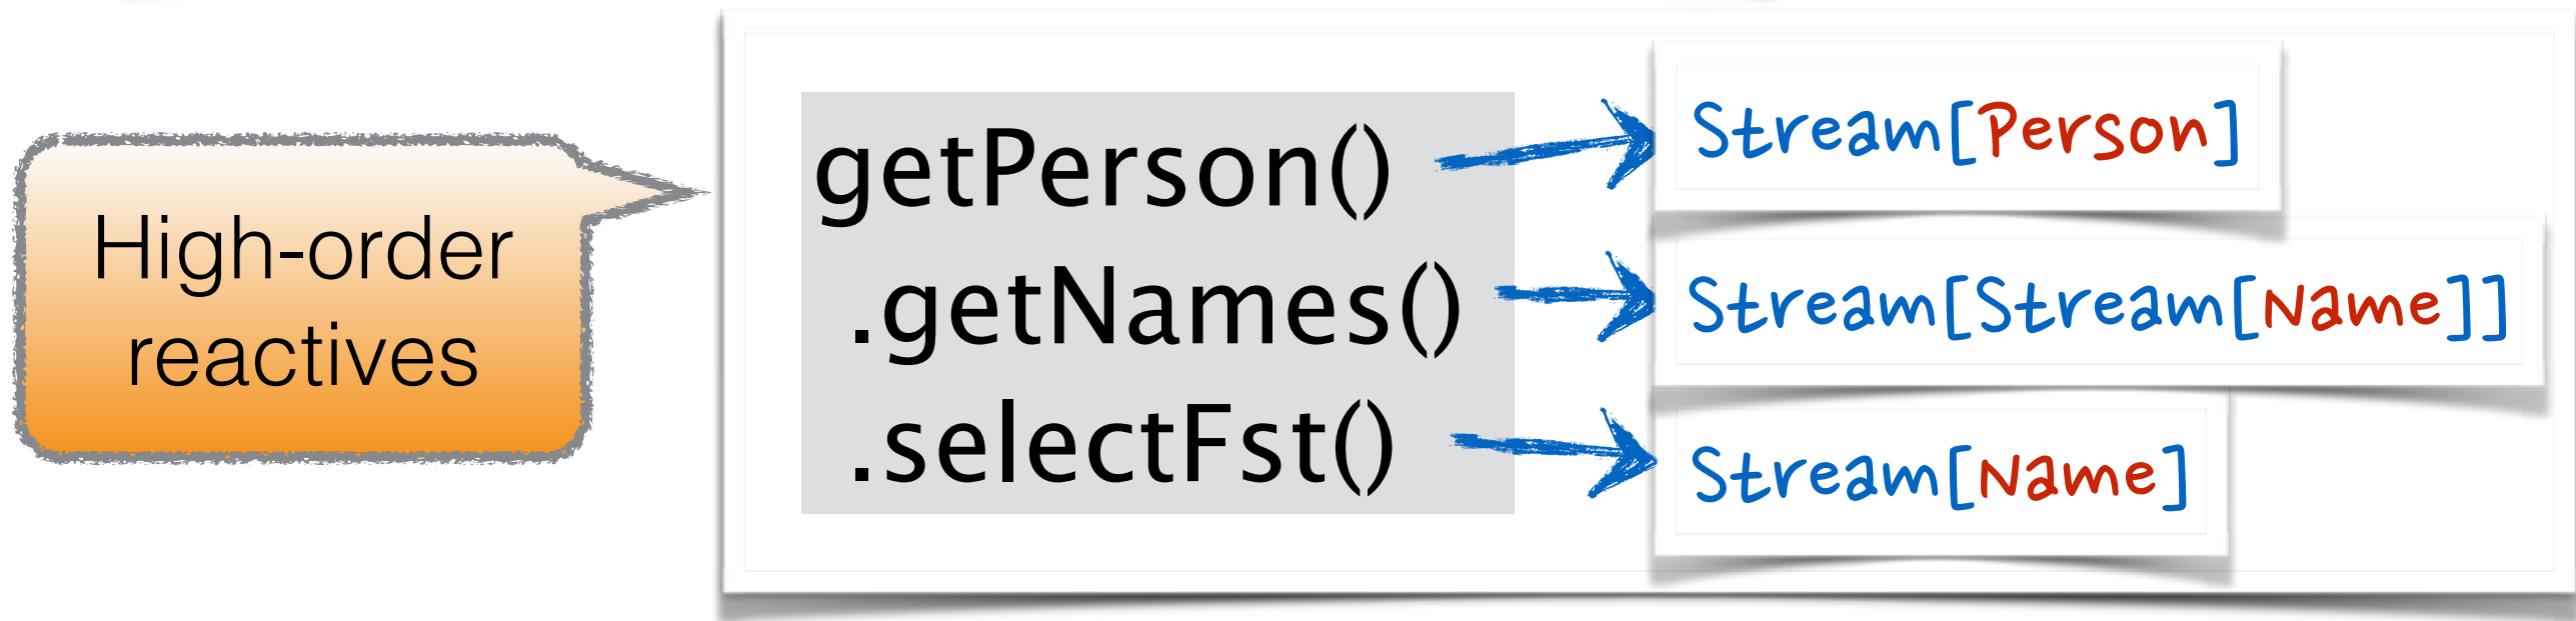
NEW: send & store more data

Evaluation

Dynamic dependencies

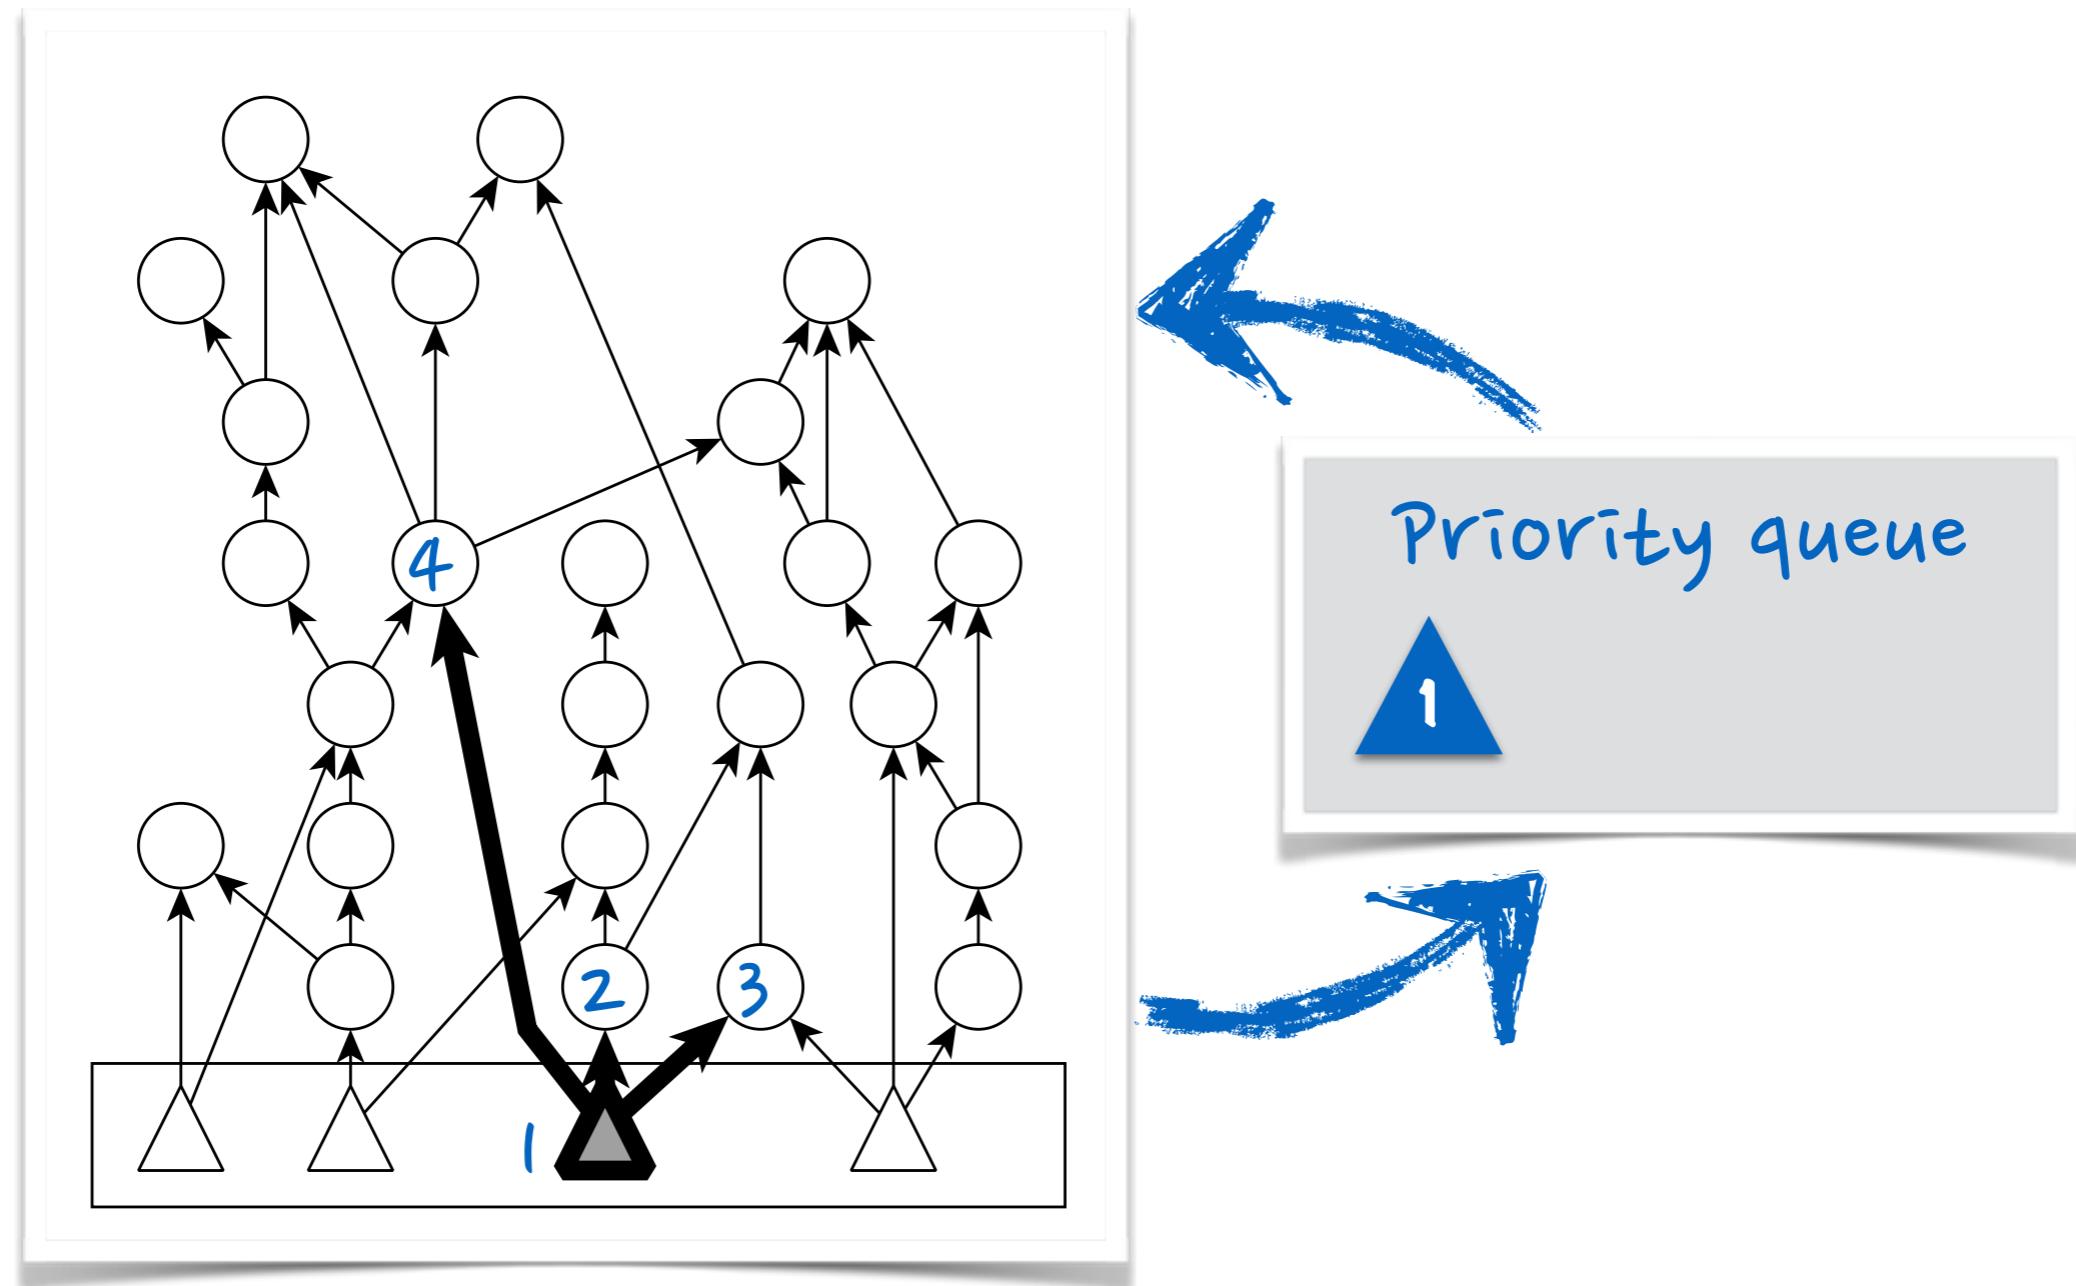


Conditionals



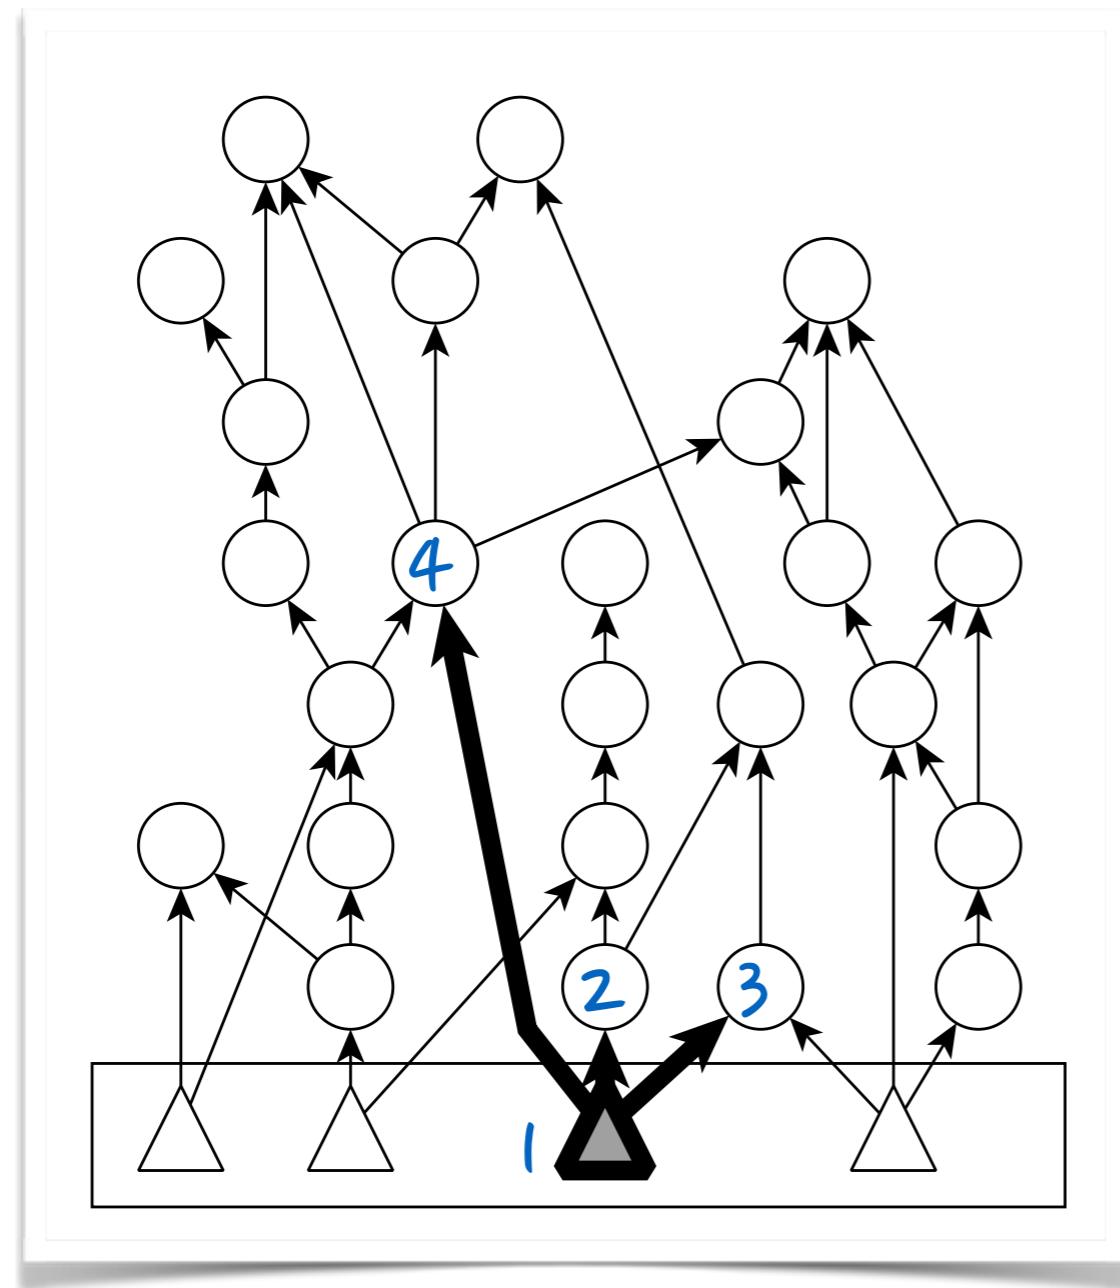
Scala.React

Topological sorting + priority queue



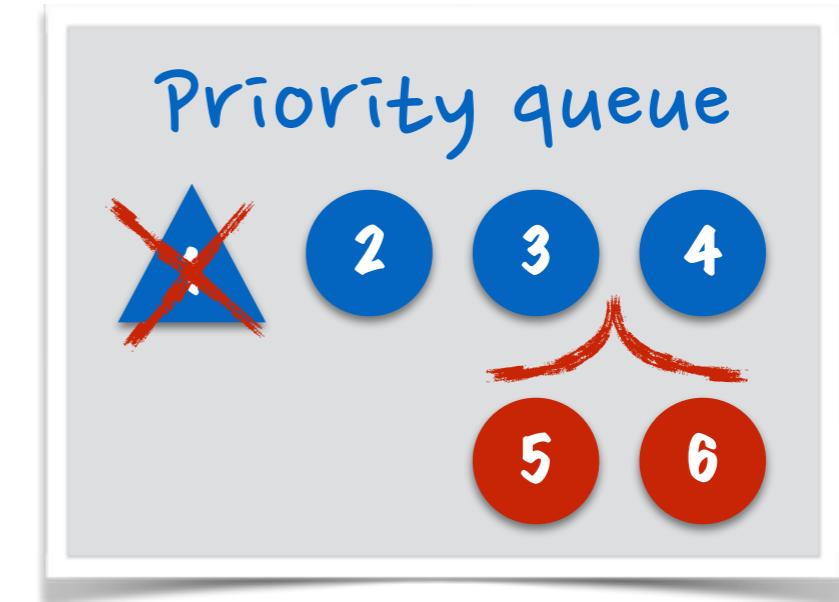
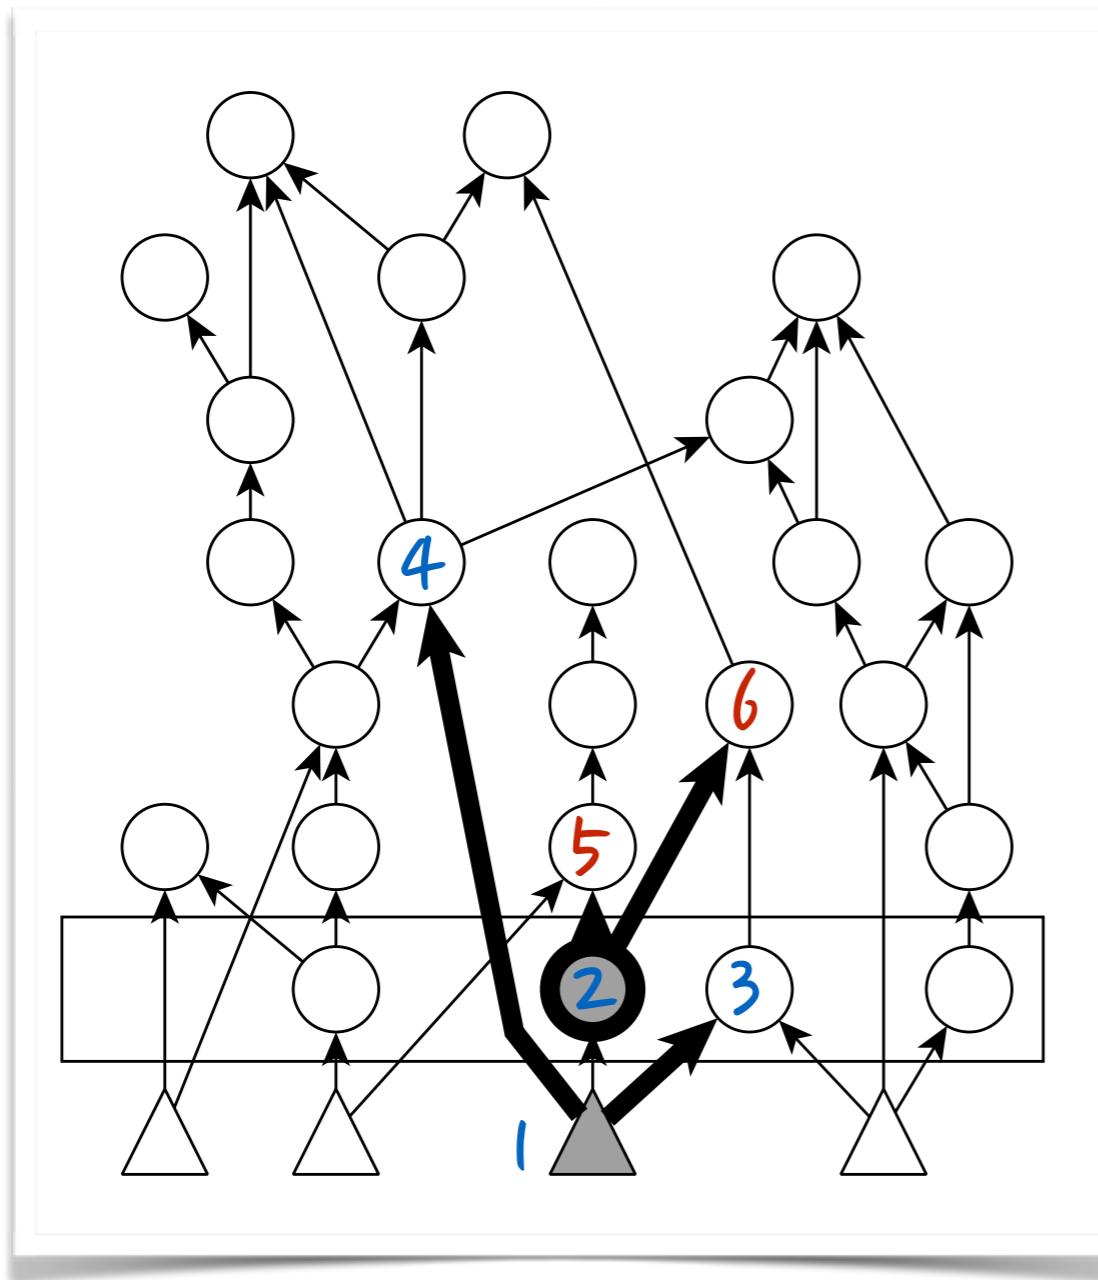
Scala.React

Topological sorting + priority queue



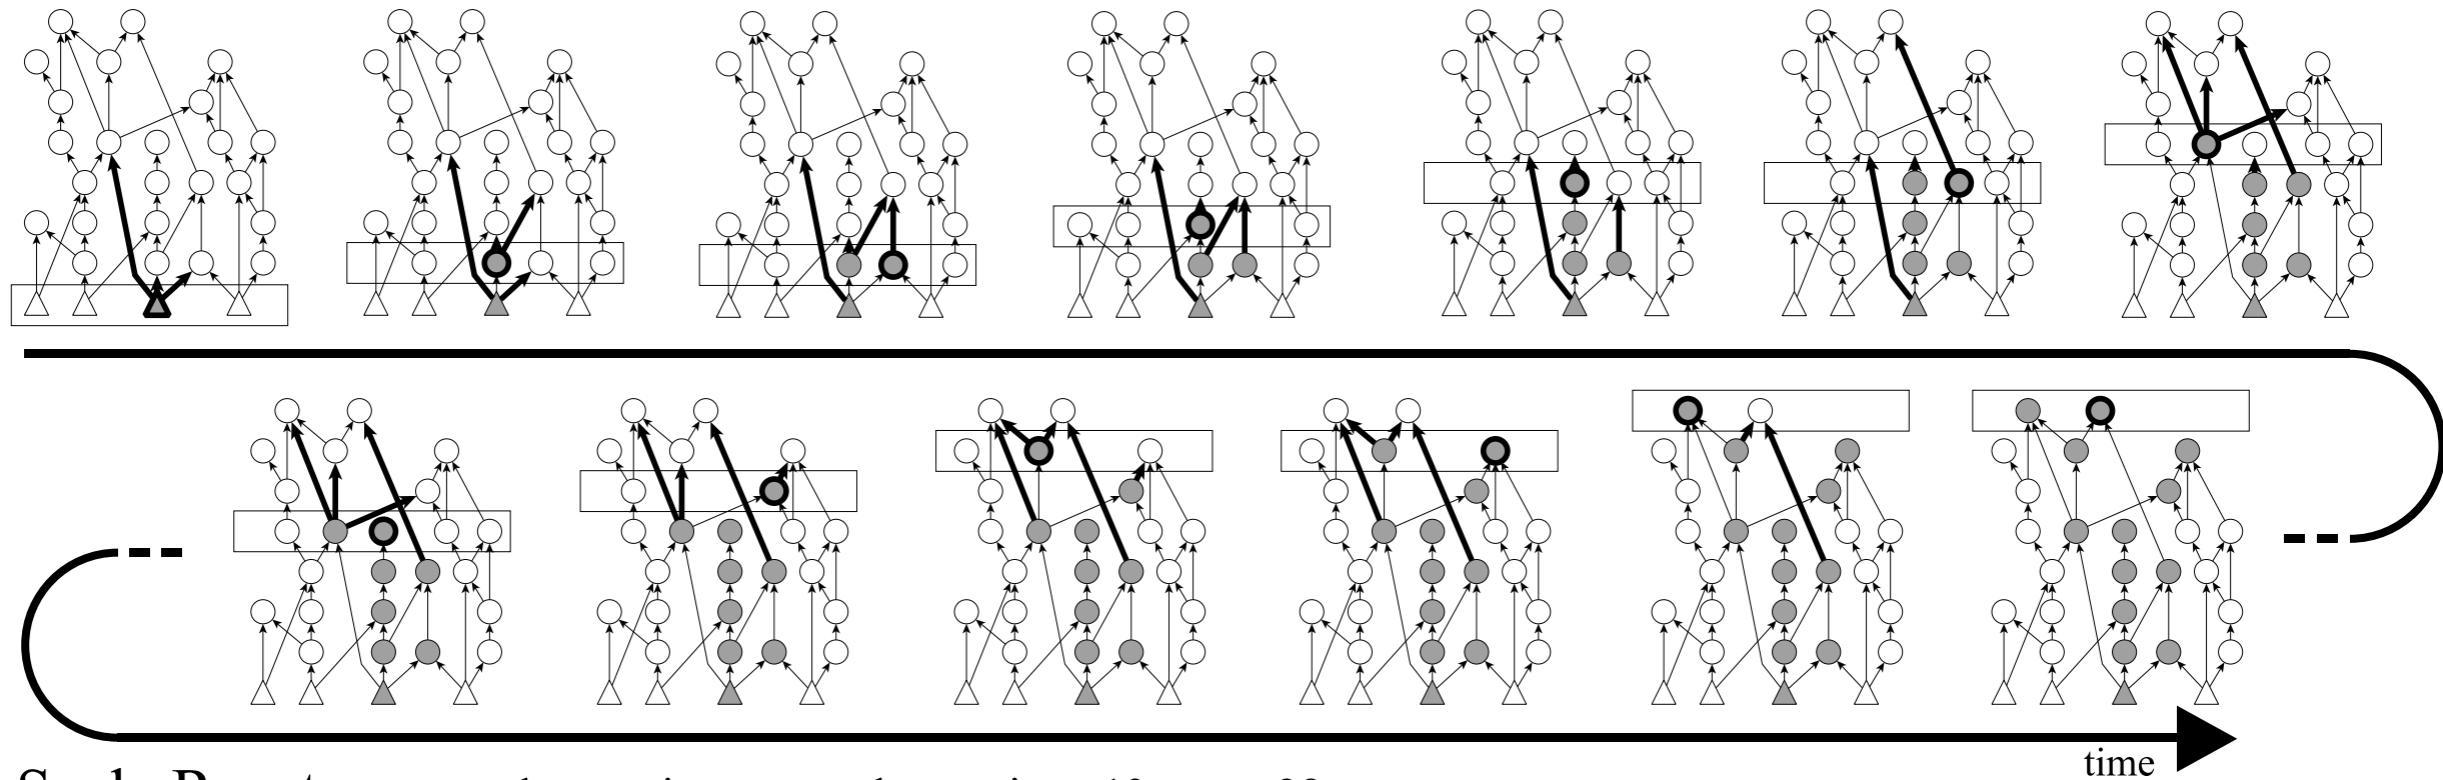
Scala.React

Topological sorting + priority queue



Scala.React

Topological sorting + priority queue



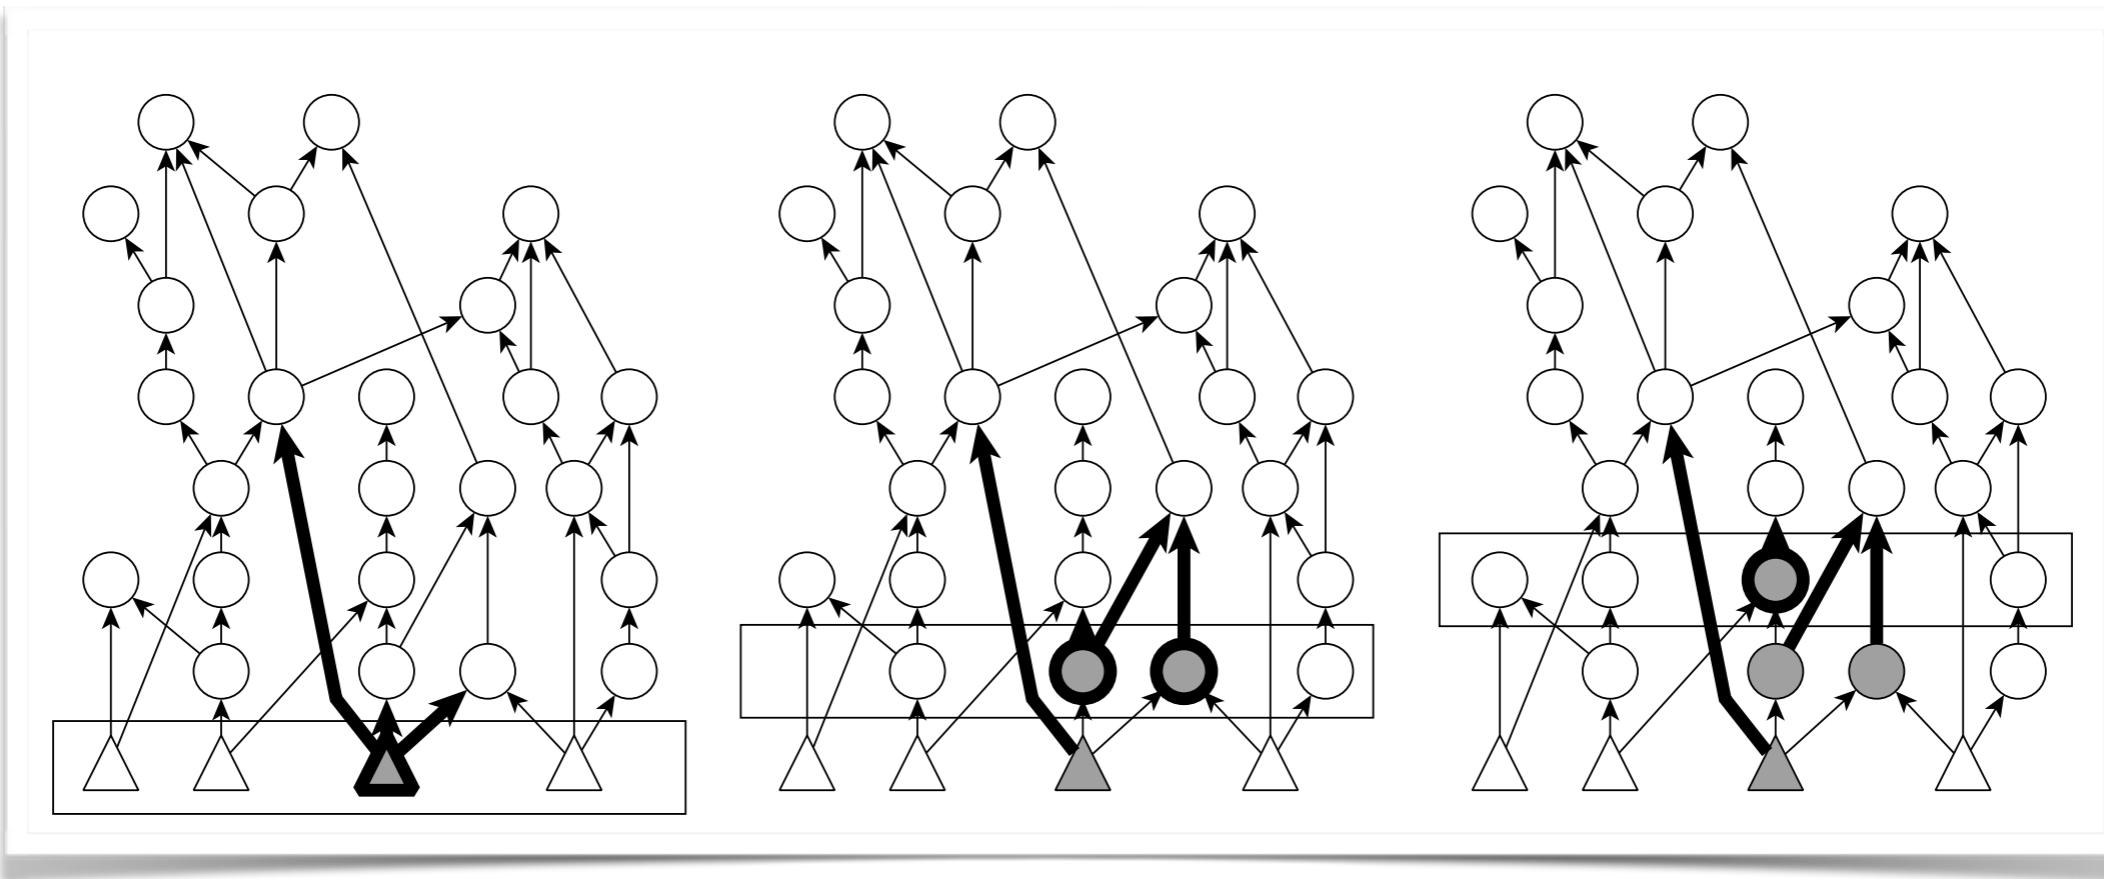
Scala.React: process layer-wise, one node at a time. 13 steps. 28 messages.

13 steps

28 messages

Scala.Rx

Parallel propagator version

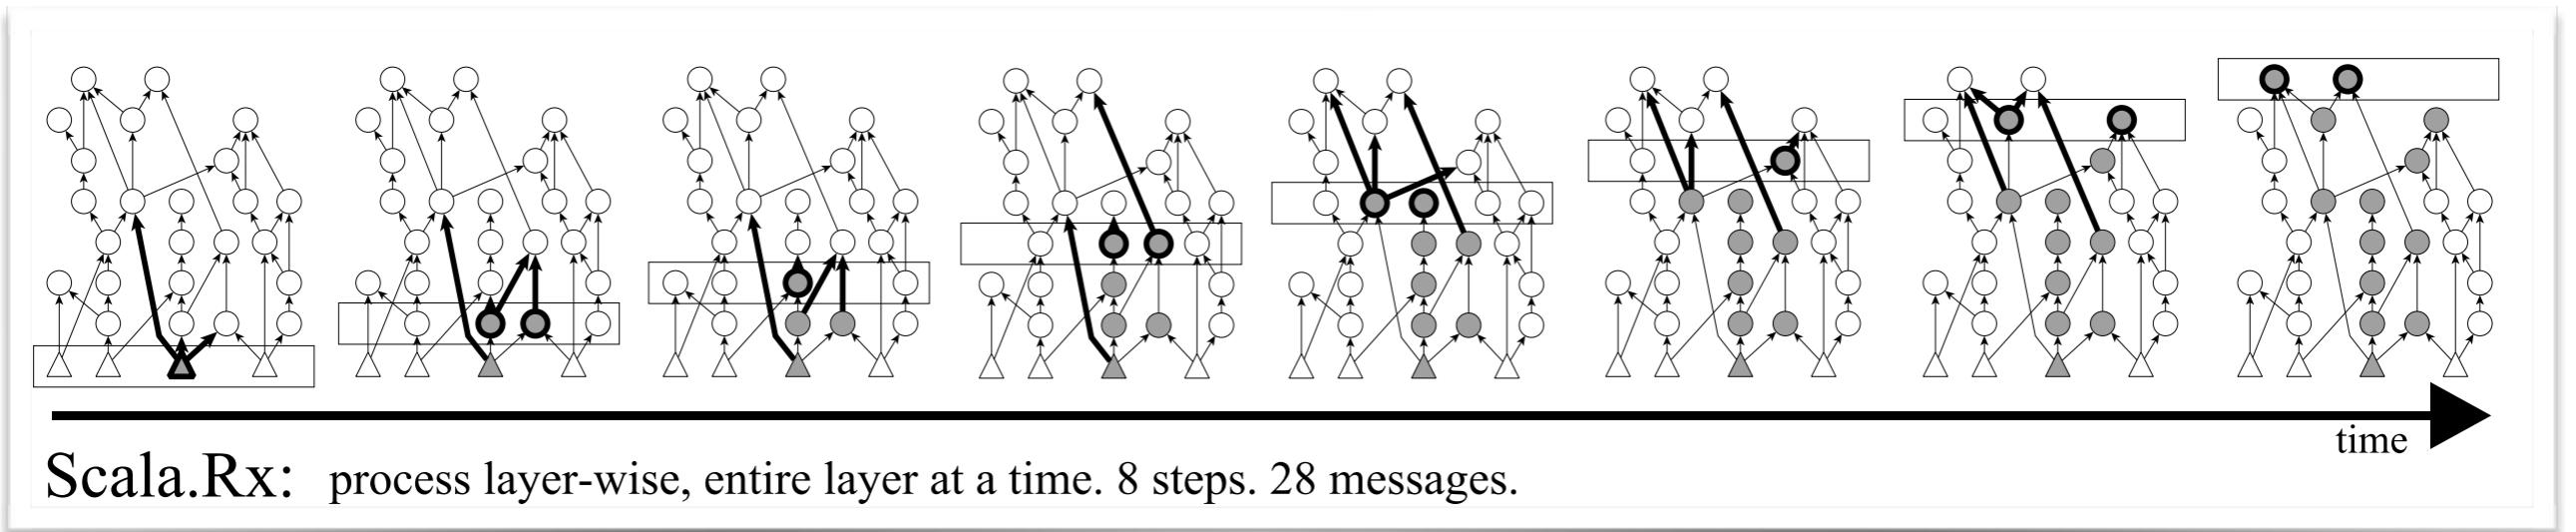


Entire layer at a time!



Scala.Rx

Parallel propagator version



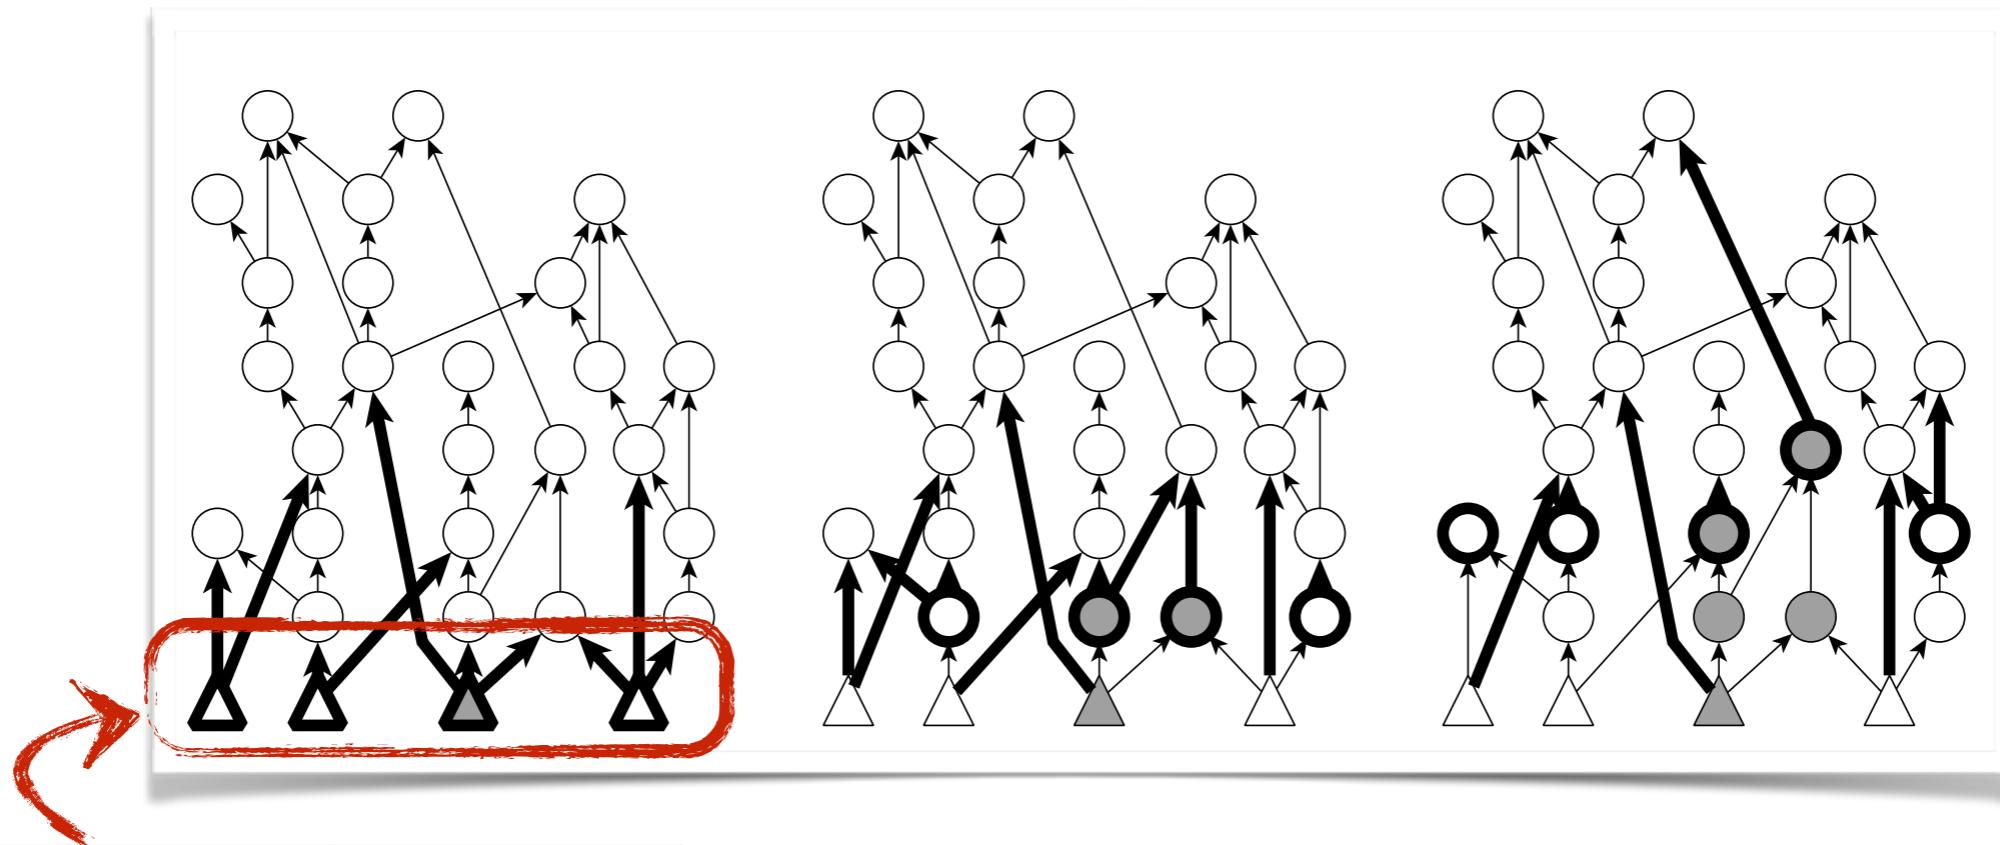
8 steps

~~13 steps~~

28 messages

ELM

Decentralised flooding

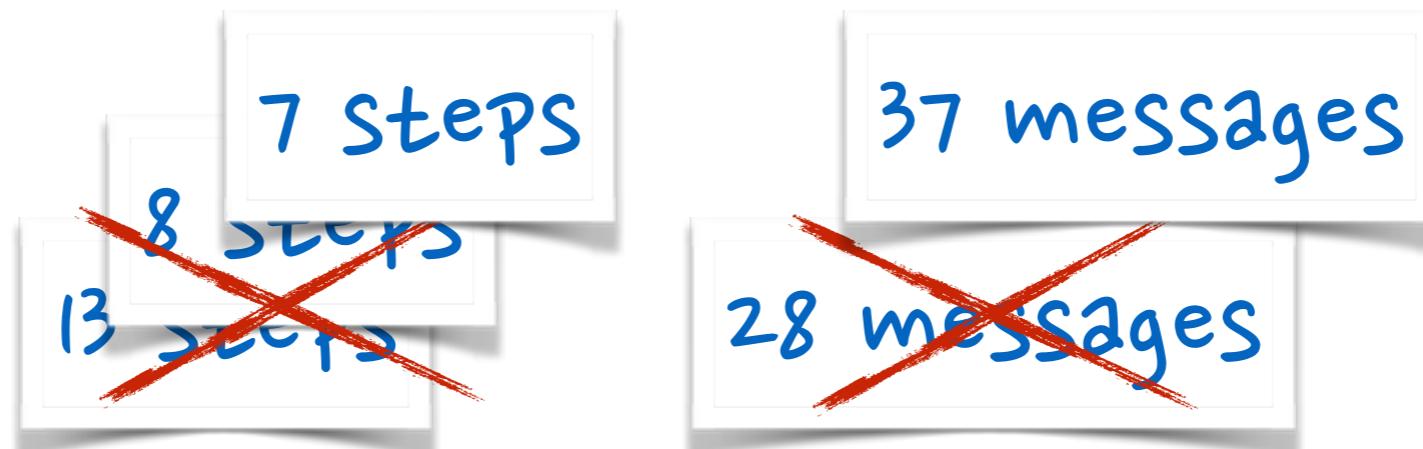
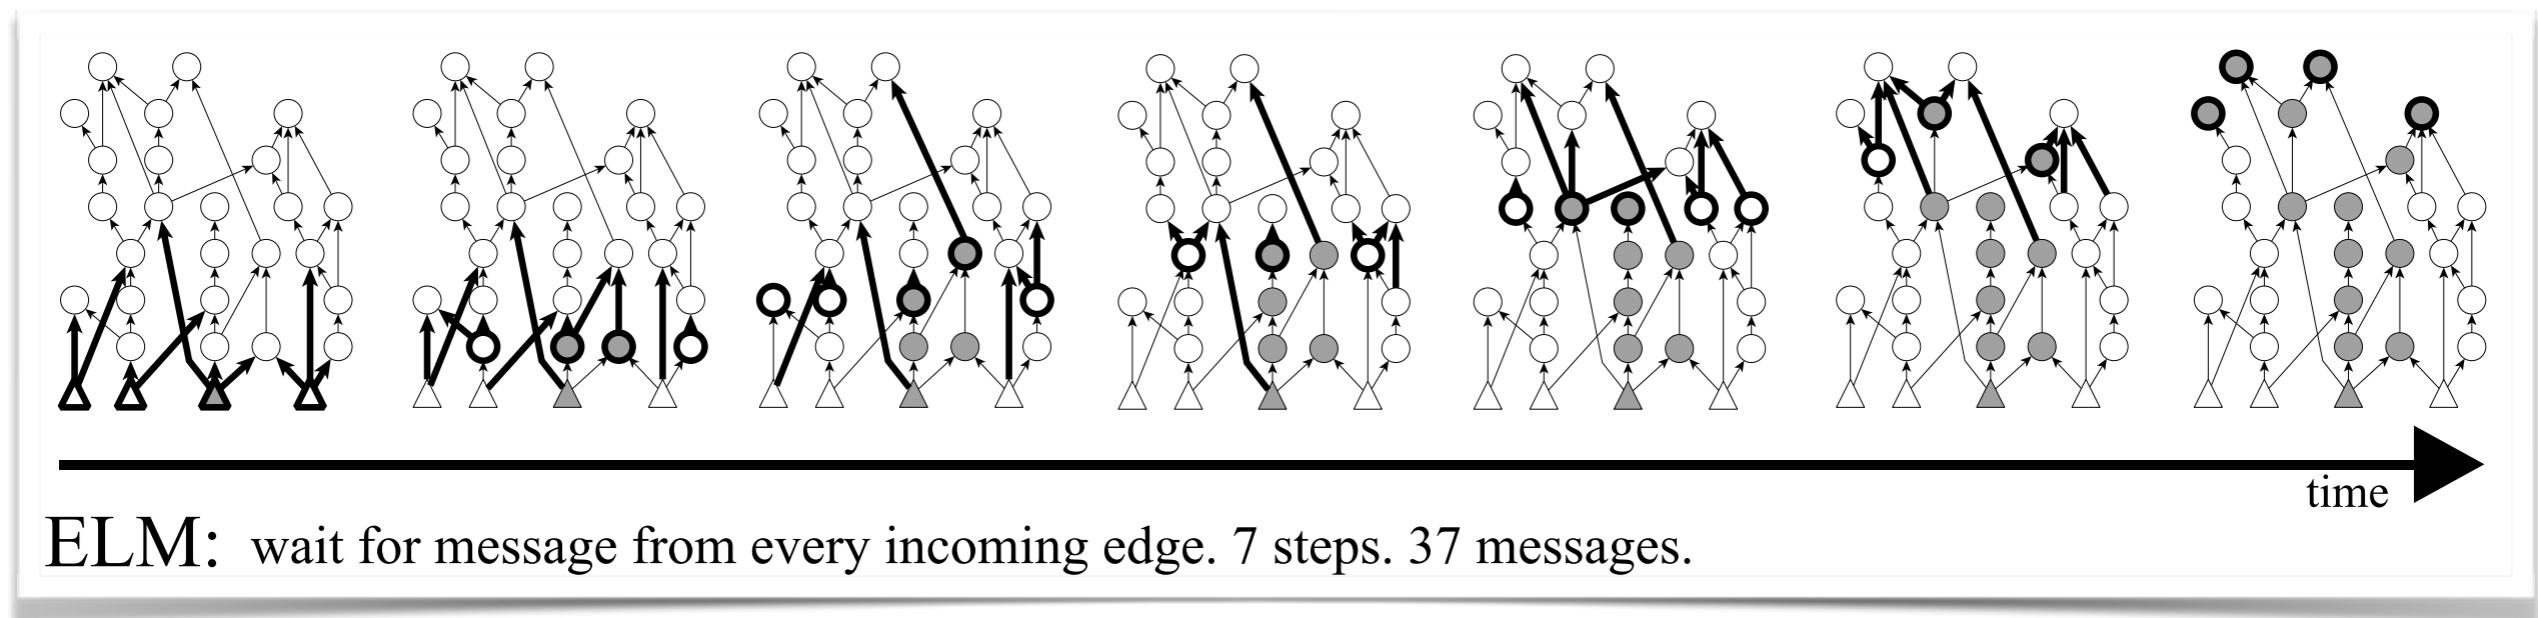


Round initiator

Each node sends:
- “change” pulse, or
- “no-change” pulse

ELM

Decentralised flooding



ELM

Decentralised flooding

Supports pipelining

BUT

No high-order reactives

ELM^S

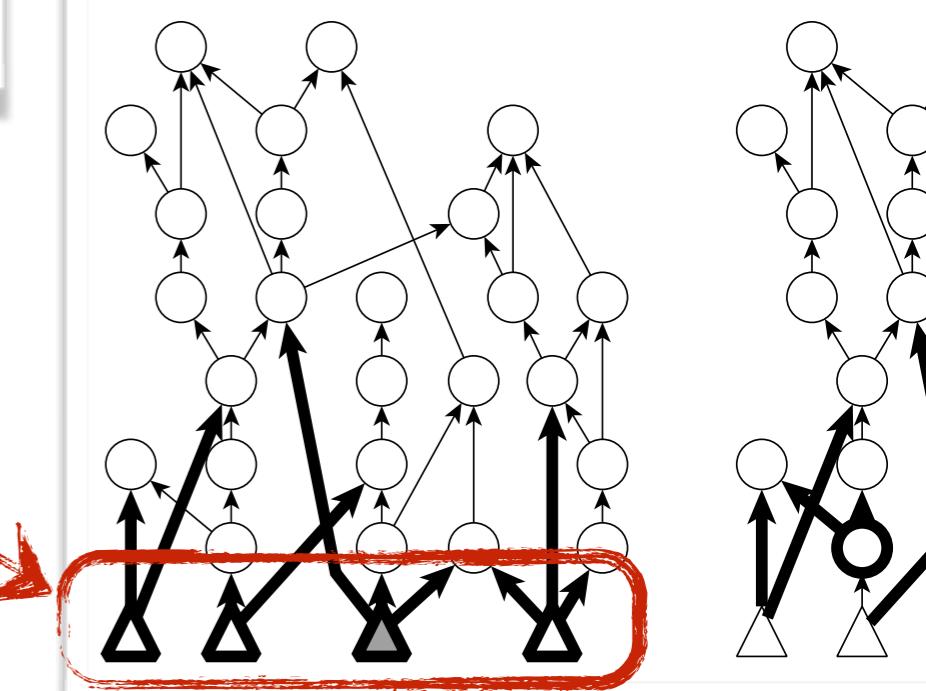
Supports HOR

BUT

No pipelining

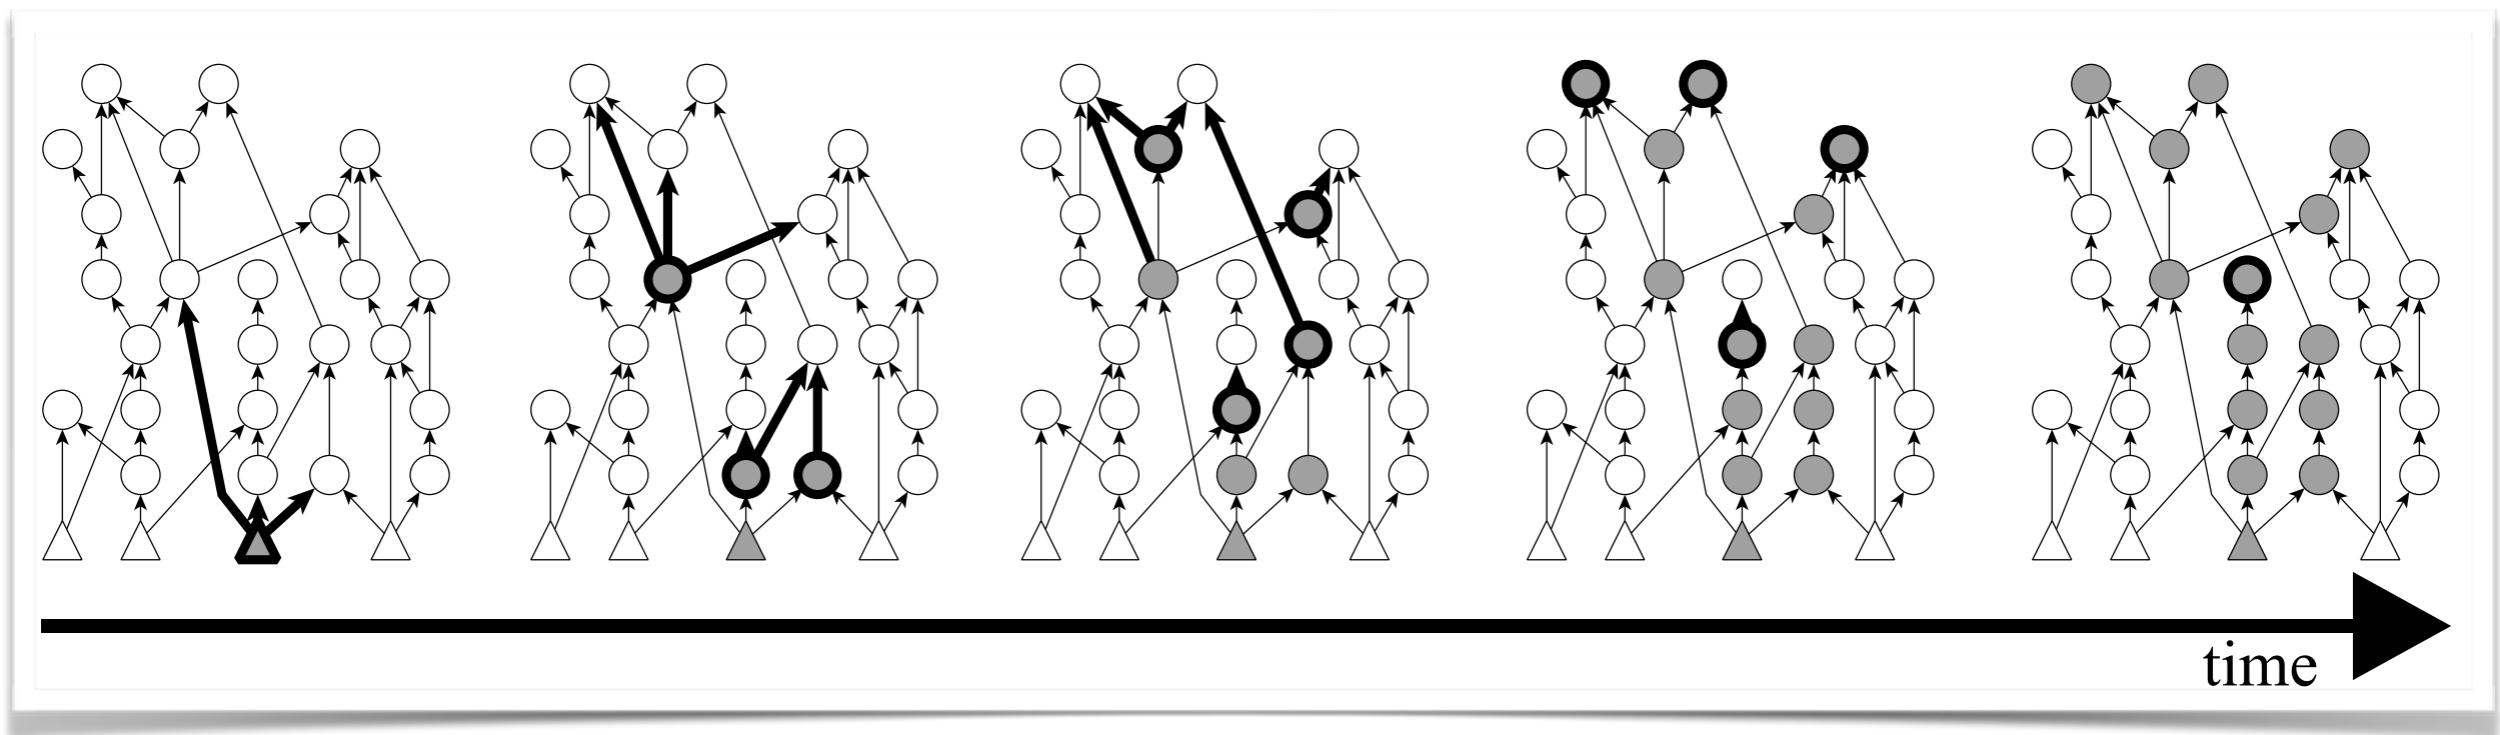
Round initiator

blocks until
the end of
the round

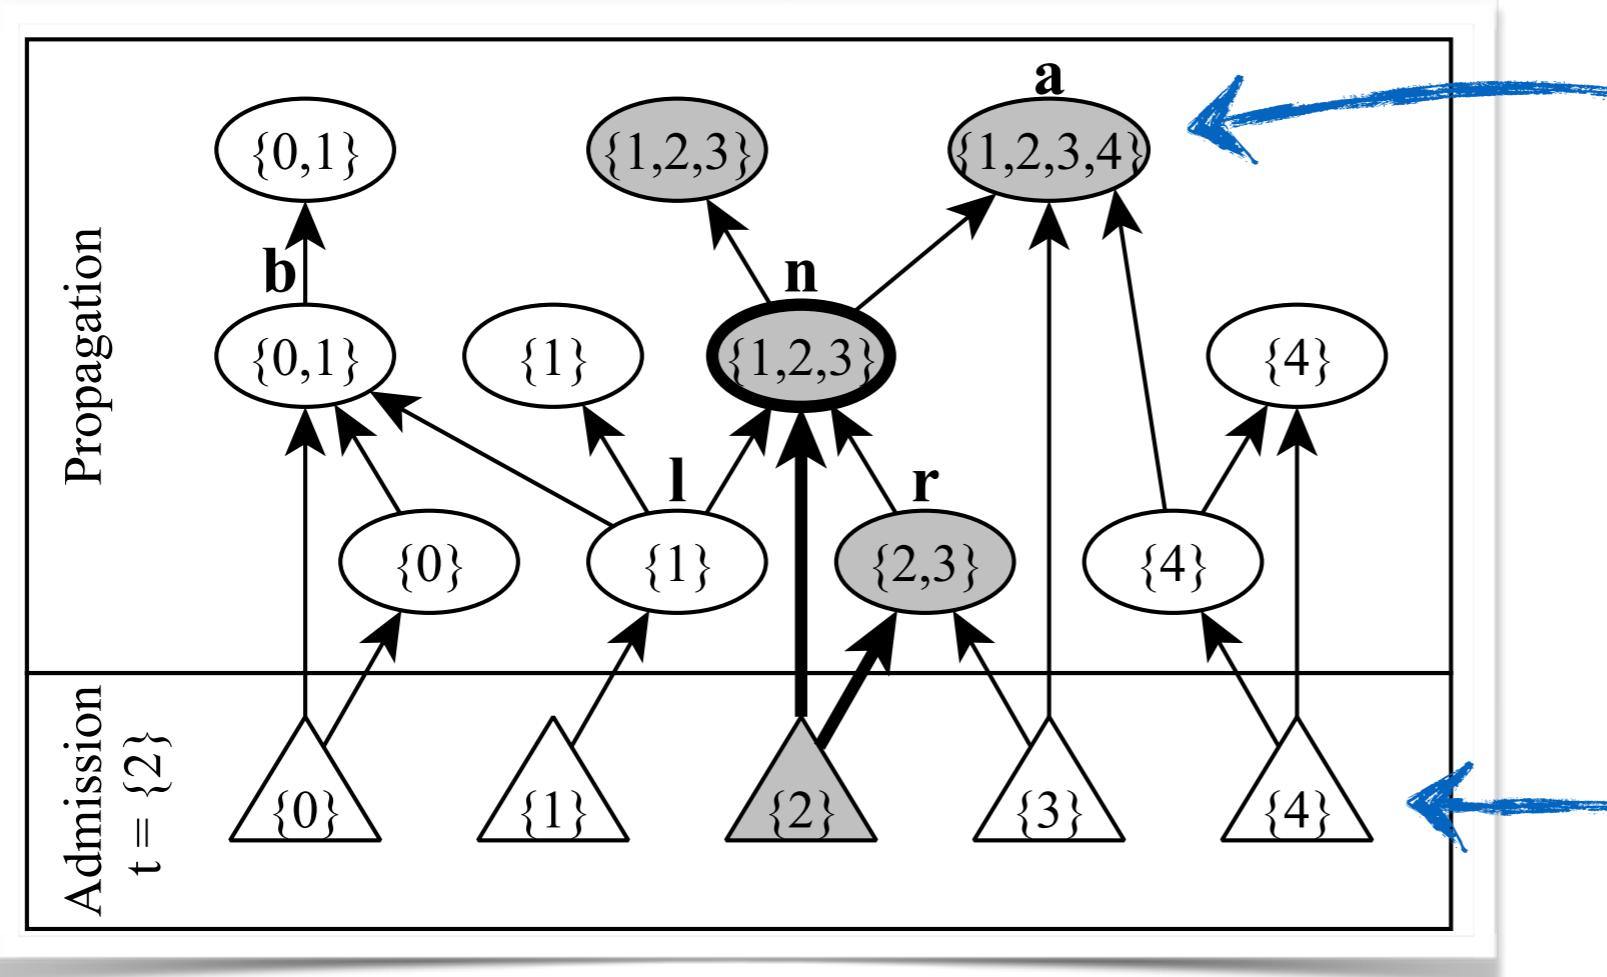


SID-UP

New algorithm



SID-UP



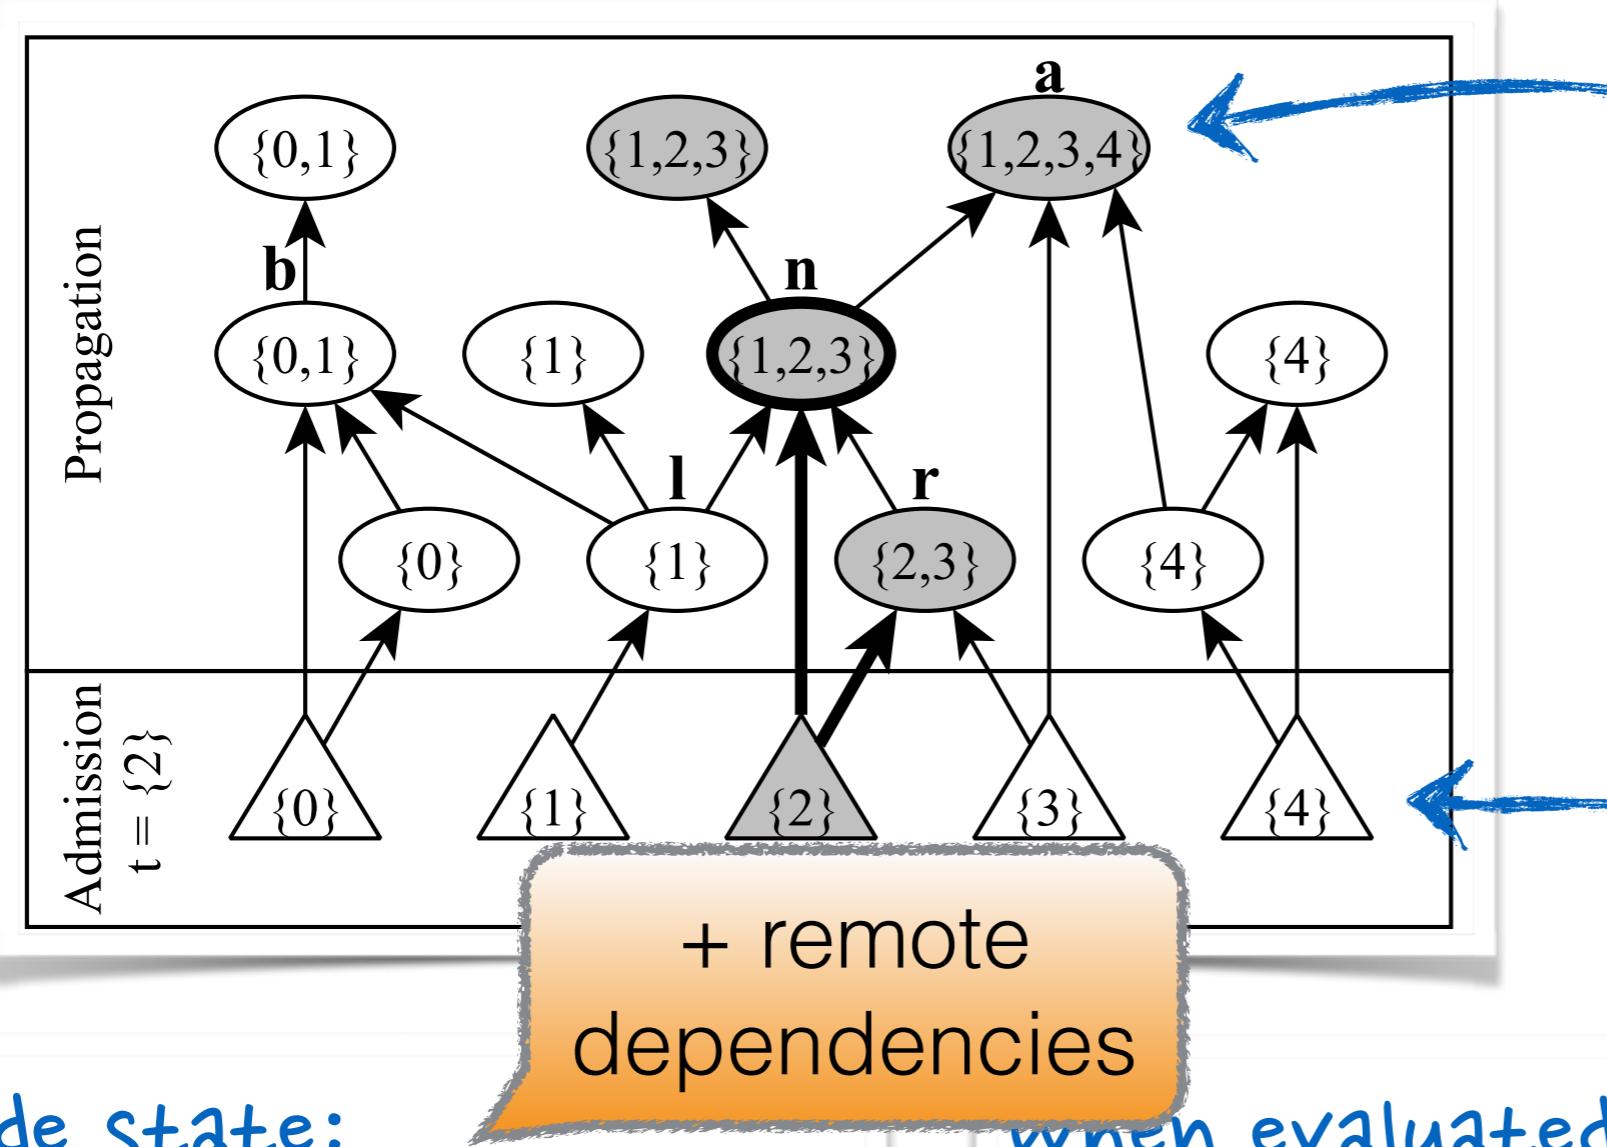
Node states:

pending (not pulsed),
changed, or unchanged

When evaluated, each node sends:

- changed + changed sources, or
- unchanged + changed sources

SID-UP



Node state:

pending (not pulsed),
changed, or unchanged

when evaluated, each node sends:

- changed + changed sources, or
- unchanged + changed sources

know
dependent
sources

unique ID's

SID-UP

glitch-freedom

No unbounded **one-to-many** communication

No coordinator when propagating

Faster than Scala.*
As fast as ELM

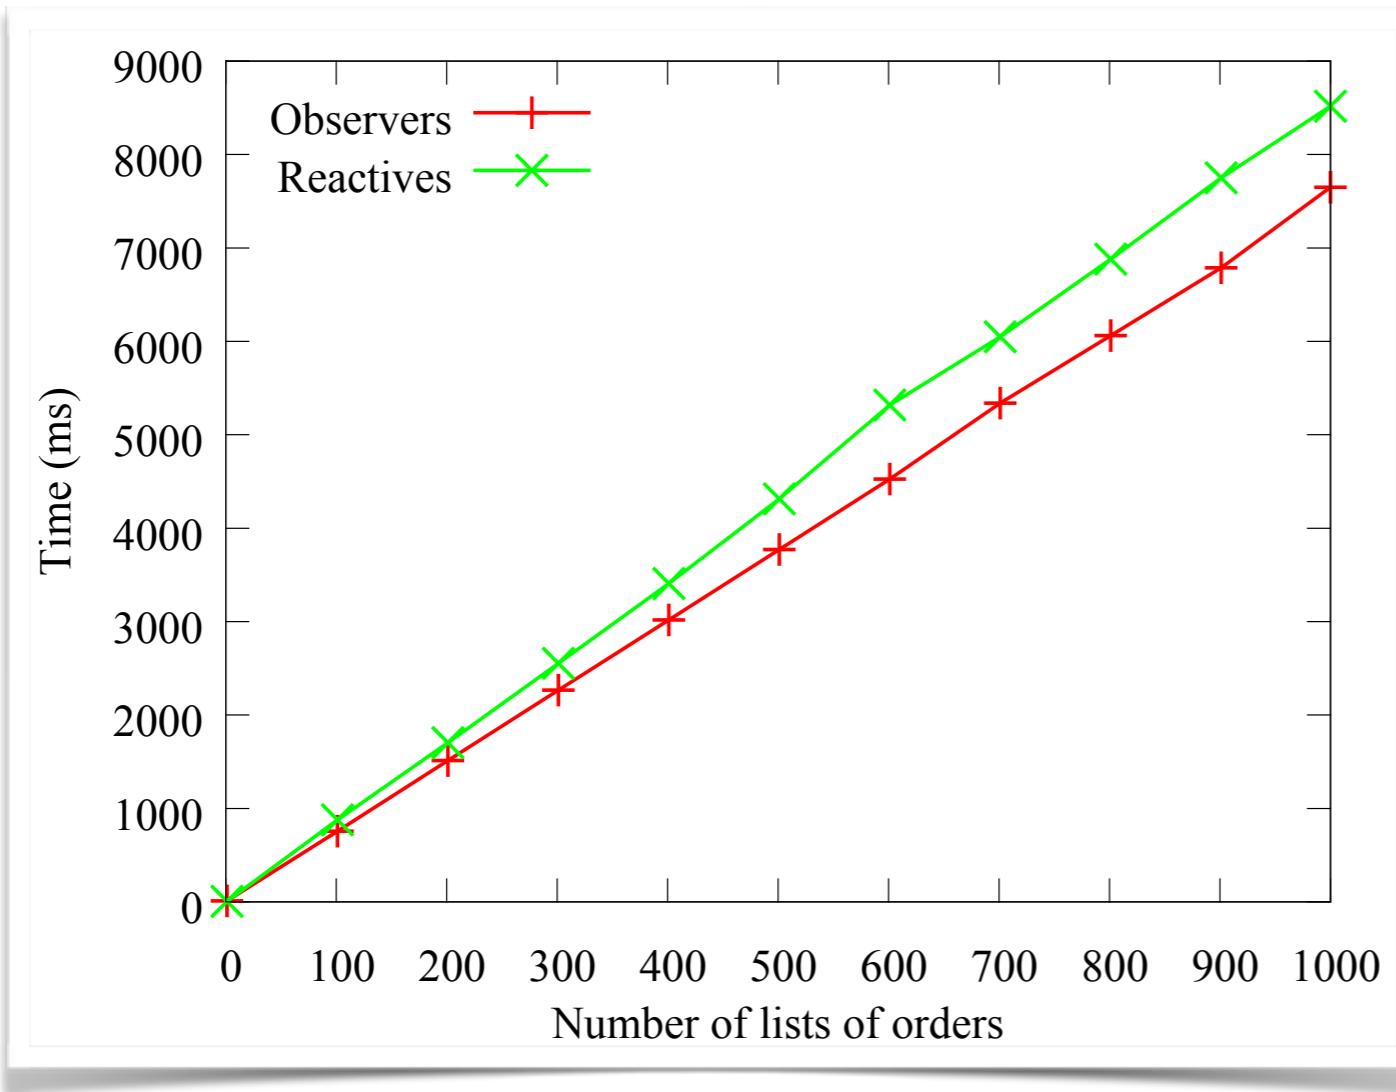
Less messages

Support for dynamic dependencies

No pipelining

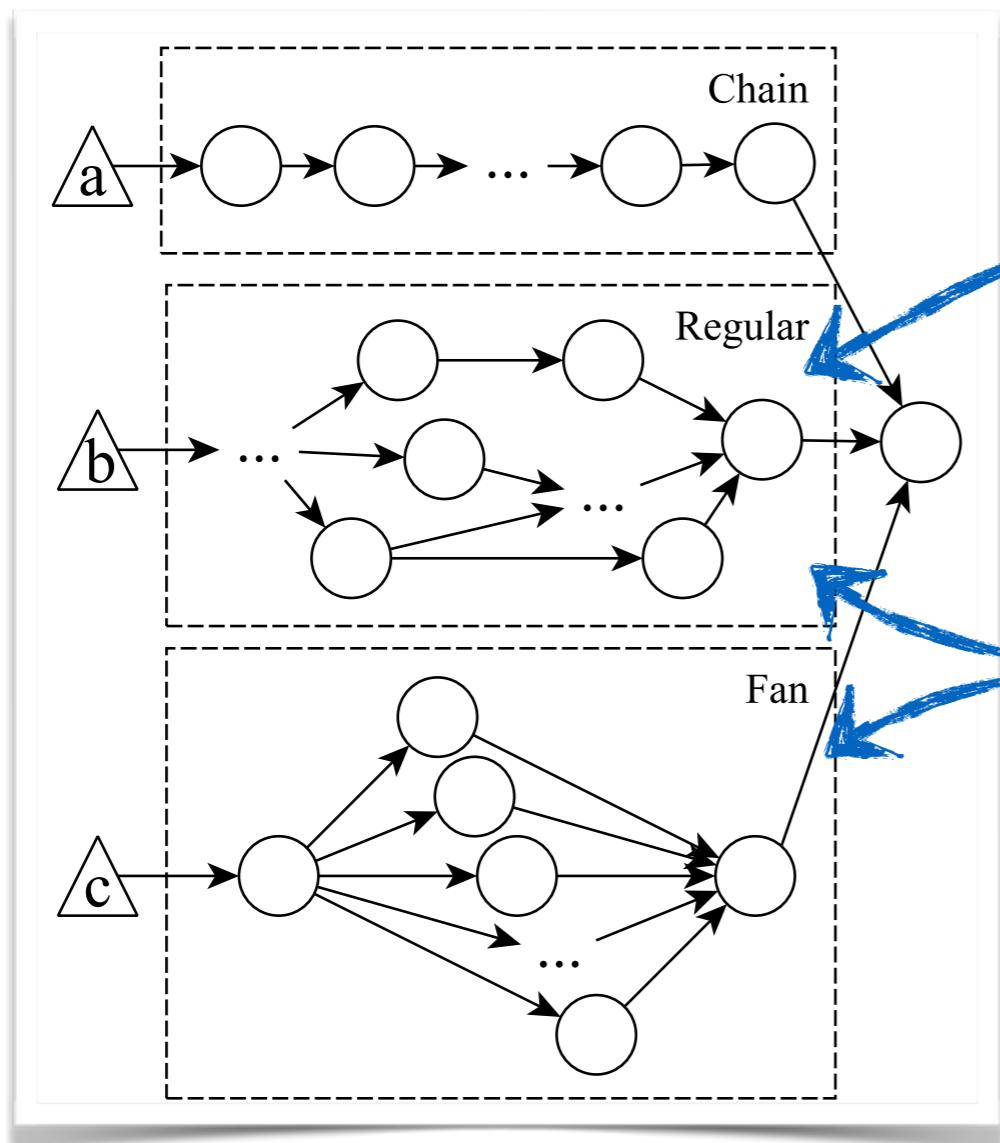
Finite & acyclic dep. graphs

Performance cost



only $\pm 10\%$ slower
than observers +
manual glitch control

Benchmarks



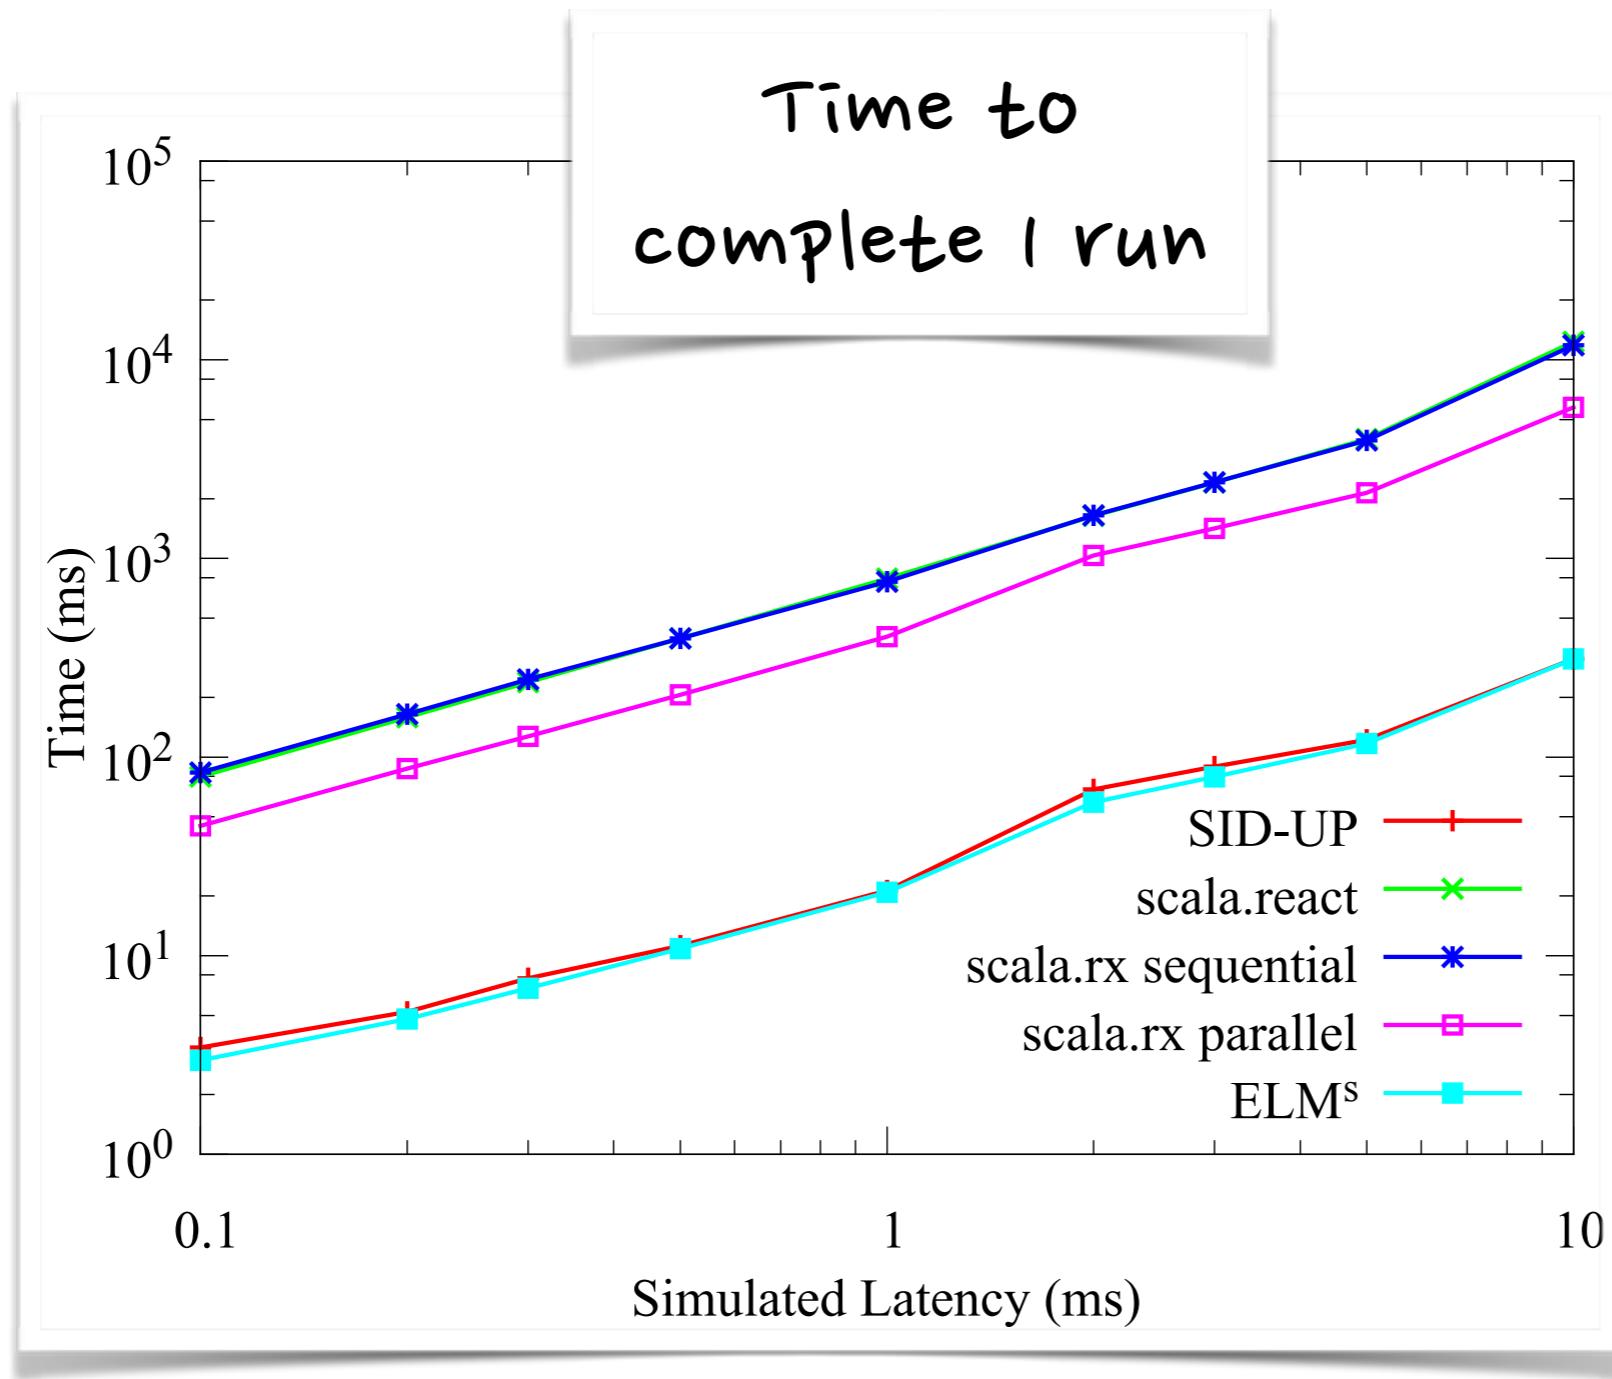
based on 20 previous
case-studies

some nodes do not
produce changes

Simulated
network latency

$$3 + 25*3 + 1$$

Benchmarks

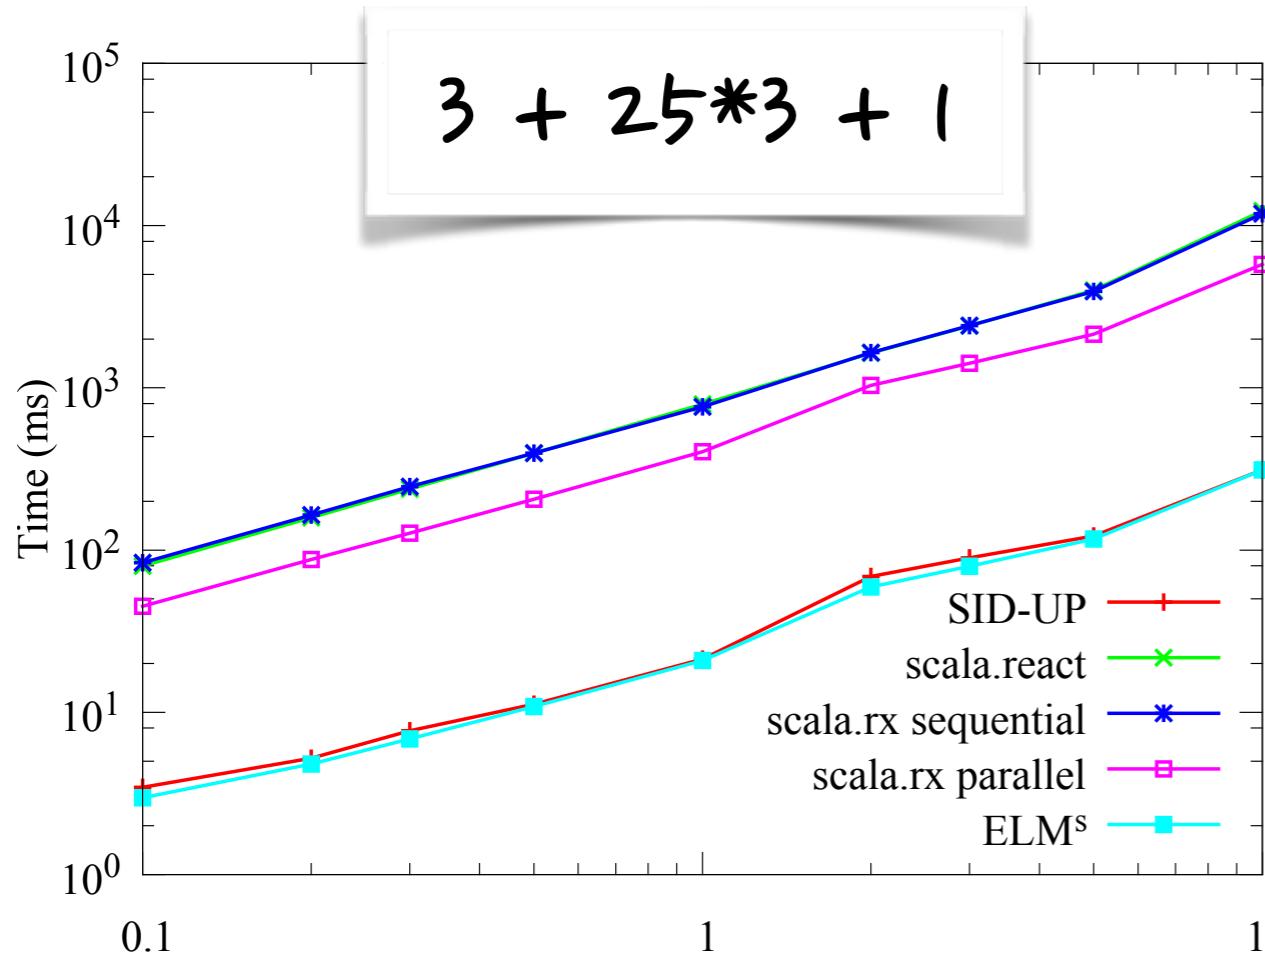


not ELM

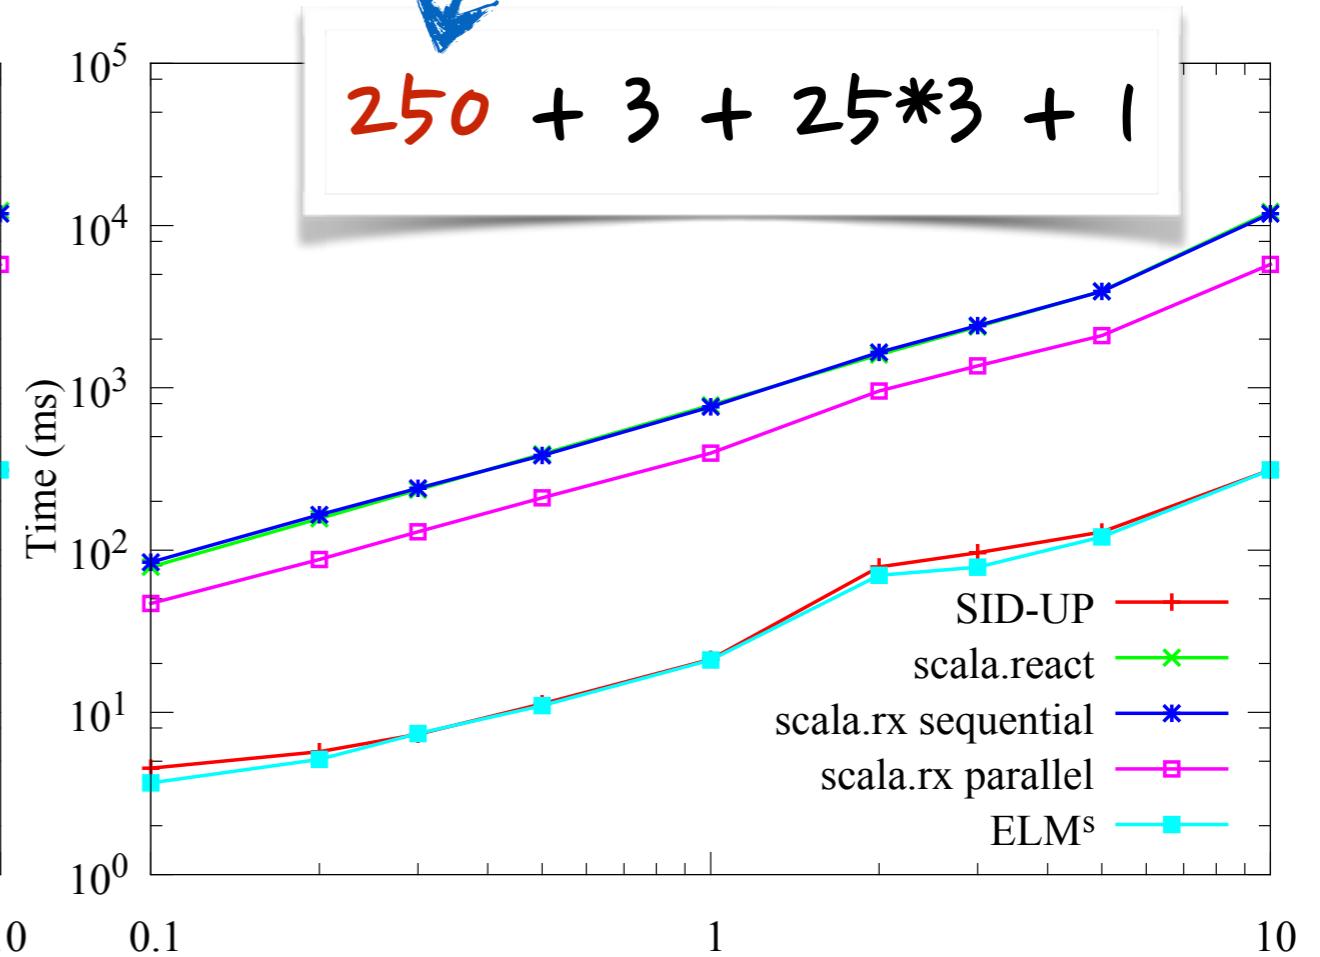
only 1 run

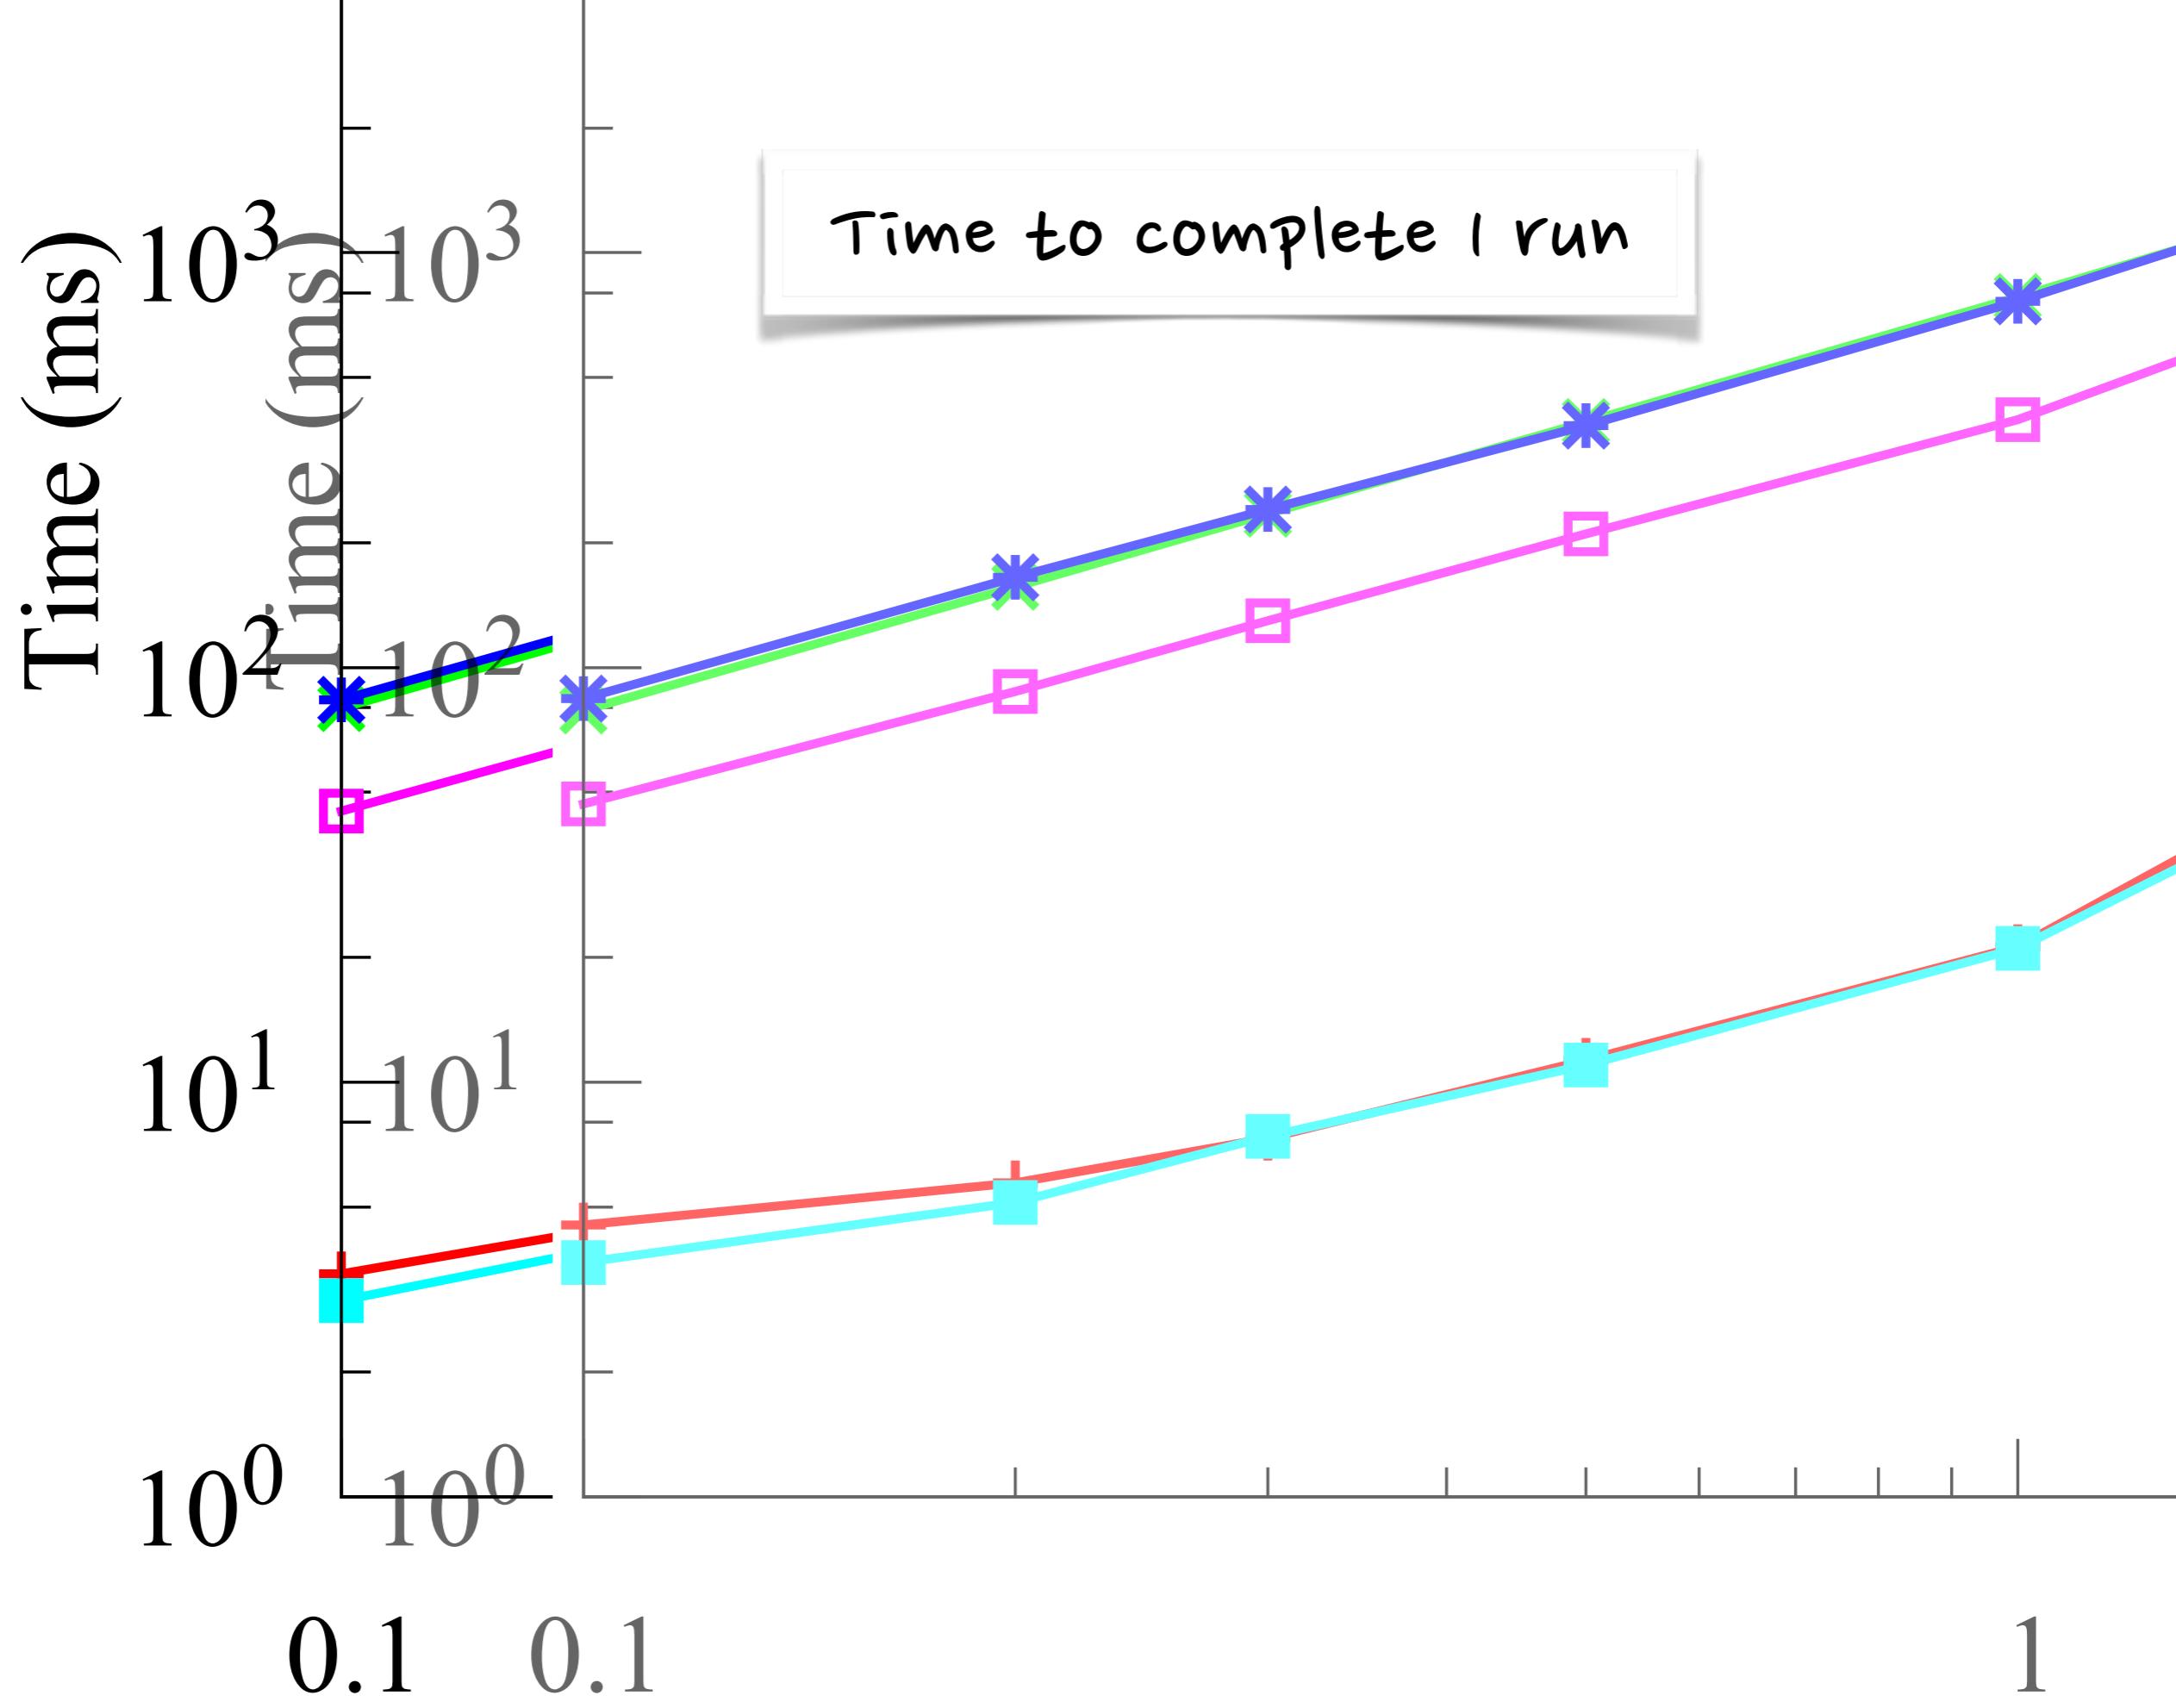
Benchmarks

Time to complete 1 run



unchanged sources

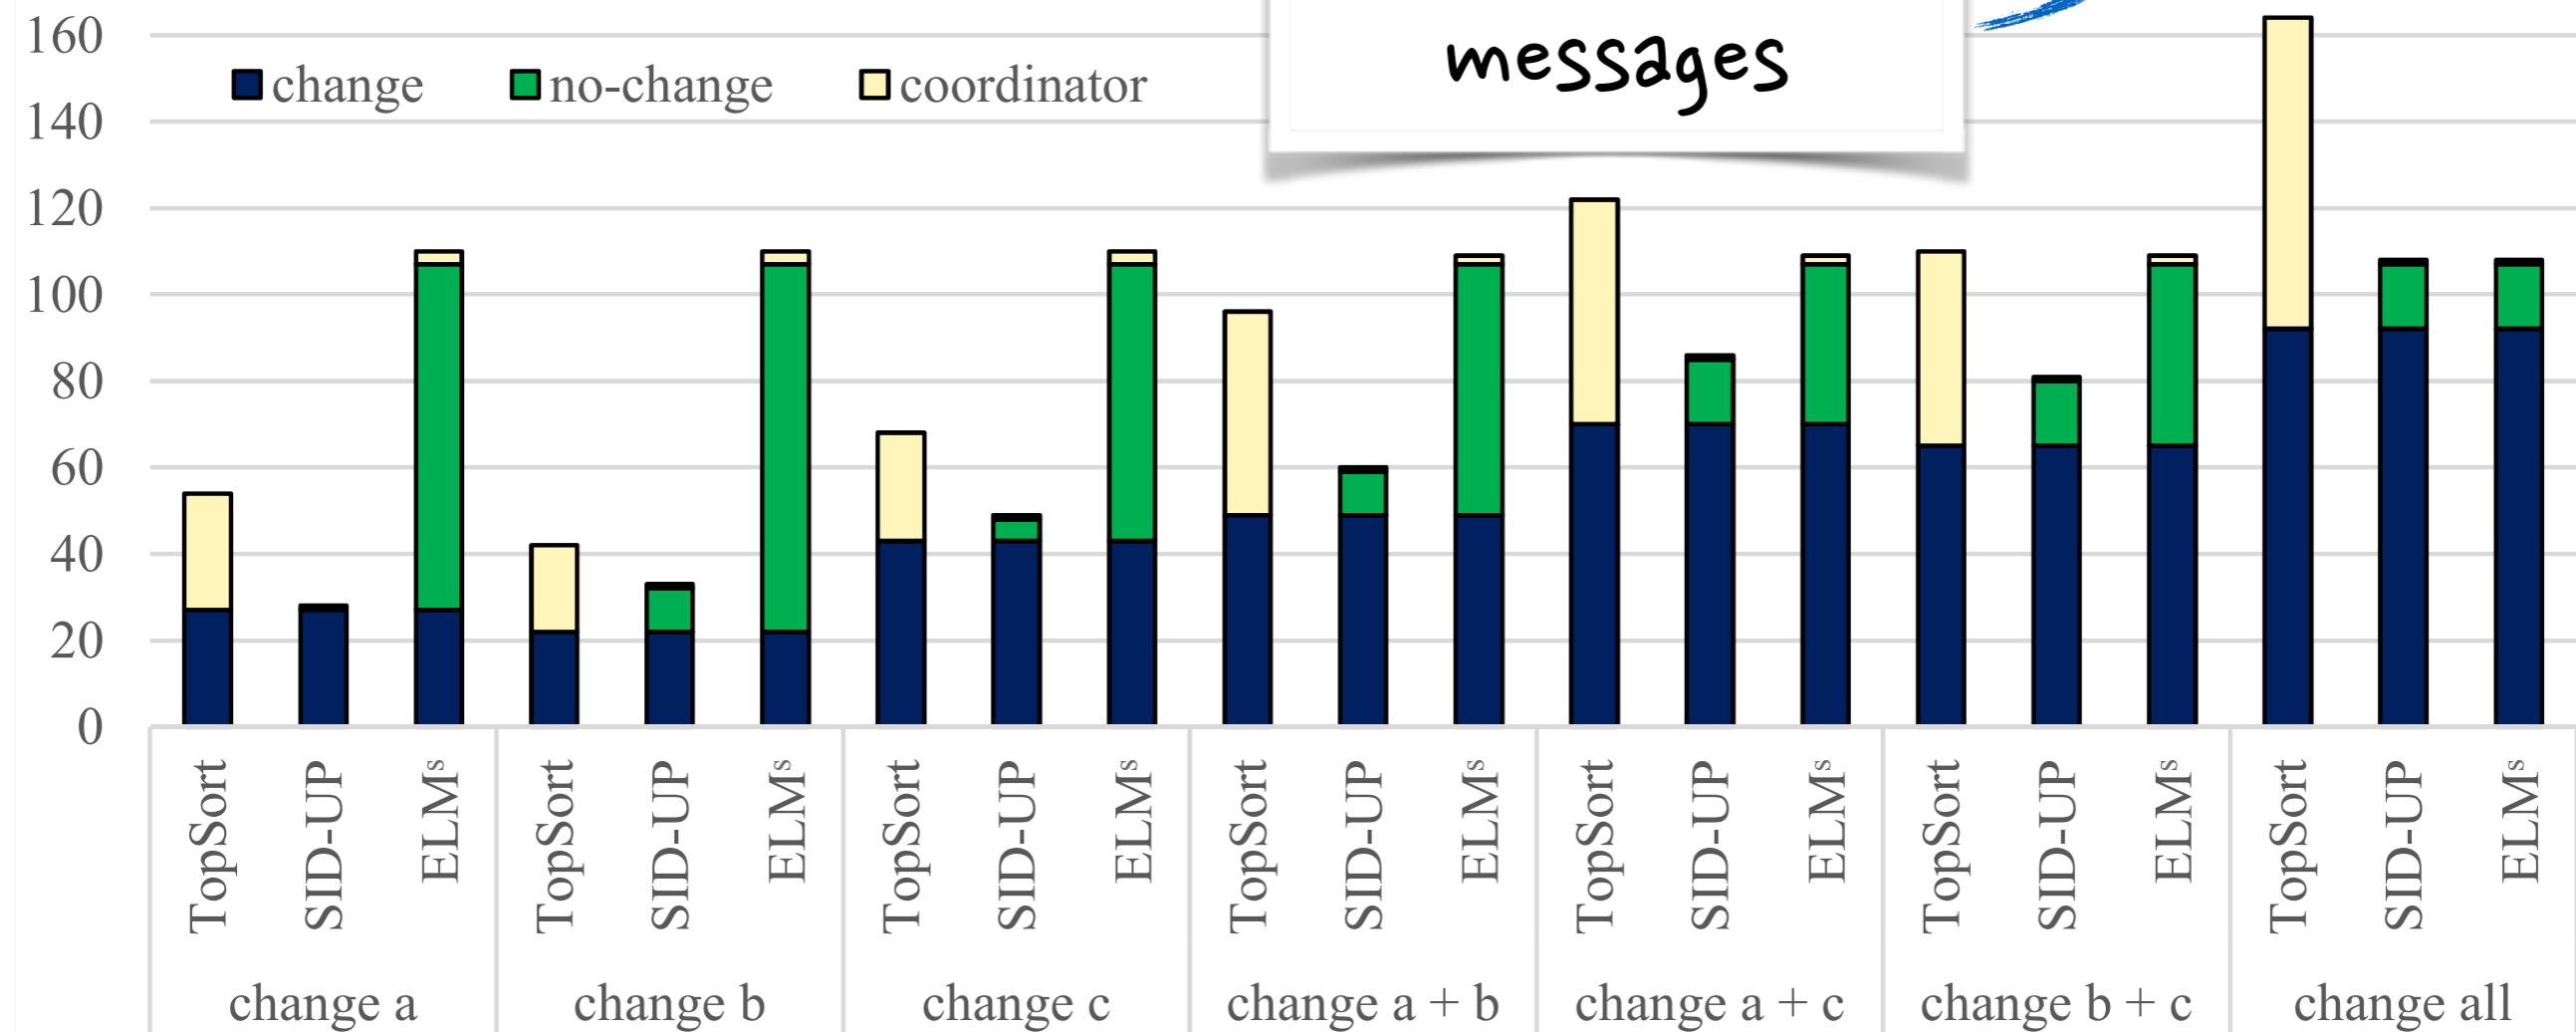




Benchmarks

to show it is better than ELM^s

Number of messages



Some discussion

ELM: Pipelining support - not evaluated (only 1 run)

Scala.*: Pipelining possible?

ELMs^s + SID-UP: How to detect the end of a run?

Lots of changed sources: not tested

Scalability? (lock at each round)

Dynamic dependencies:

can affect topology? Assumptions?

distributed
clock?

Optimisations

REBLS '14 - Splash workshop

“Optimizing Distributed REScala”

Joscha Drechsler and Guido Salvaneschi

If has 1 incoming dep.:

avoid iterations,
intersections, waiting

If has no dyn. dep.:

just **count** incoming
pulses

include **valueChanged** and
sourcesChanged in pulse

Optimisations

REBLS '14 - Splash workshop

"Optimizing..."

Joscha D.

Benchmarks without
injected network delays

If has 1 incoming dep.:

avoid iterations,
intersections, waiting

If has no dyn. dep.:

just **count** incoming
pulses

include **valueChanged** and
sourcesChanged in pulse