

We Have a DREAM: Distributed Reactive Programming with Consistency Guarantees

Alessandro Margara, Guido Salvaneschi

Presented by Wilfried Daniels

Introduction

- Designing, implementing and maintaining reactive systems is difficult
 - Asynchronous callbacks
 - Hard to trace/understand control flow

→ Solution: Reactive Programming

Introduction

- Key concepts:
 - time-varying values
 - tracking of dependencies
 - automatic propagation of changes

```
1 var a: int = 10  
2 var b: int = a + 2  
3 println(b) // 12  
4 a = 11  
5 println(b) // 12
```

Imperative

```
1 var a: int = 10  
2 var b: int := a + 2  
3 println(b) // 12  
4 a = 11  
5 println(b) // 13
```

Reactive

Introduction

- Advantages vs. classic event-based arch:
 - No explicit update logic
 - Declarative specification of dependencies
 - Runtime manages correct propagation (e.g. glitch freeness/consistency)
- This work focuses on distributed reactive programming (DRP)


Introduction

- Previous DRP solutions do not guarantee distributed consistency (only local)
- This paper presents DREAM , a **Distributed REActive Middleware** with three different levels of consistency guarantees

Background and Motivation

- Motivation for different levels of consistency
- Running example: financial application system

```
1 var marketIndex = InputModule.getMarketIndex()  
2 var stockOpts = InputModule.getStockOpts()  
3 var news = InputModule.getNews()  
4  
5 // Forecasts according to different models  
6 var f1 := Model1.compute(marketIndex,stockOpts)  
7 var f2 := Model2.compute(marketIndex,stockOpts)  
8 var f3 := Model3.compute(marketIndex,news)  
9  
0 var gui := Display.show(f1,f2,f3)  
1  
2 var financialAlert := ((f1+f2+f3)/3) < MAX  
3 if (financialAlert) decrease(stockOpts)  
4  
5 var financialAlert_n := computeAlert_n(f1,f2,f3)  
6 if (financialAlert_n) adjust_n(stockOpts)
```



Background and Motivation

- Variant 1: Smartphone app
 - Just displays output of 3 models
 - No consistency required

```
var gui := Display.show(f1,f2,f3)
```

V1

- Variant 2: Models aggregator
 - Aggregates output of 3 models
 - Undertakes action when below threshold

```
var financialAlert := ((f1+f2+f3)/3) < MAX  
if (financialAlert) decrease(stockOpts)
```

V2

Background and Motivation

- Variant 2: Models aggregator
 - Requires glitch freedom
 - Assume initially **f1**:110, **f2**:95, **f3**:99 with **MAX**:100
 - New **marketIndex**: *all* models recalculate.
 - Model **f1** finishes first with **f1**: 90
→ STOCKS DECREASED (**GLITCH!**)
 - Other models finish: **f2**:111, **f3**:103

```
var financialAlert := ((f1+f2+f3)/3) < MAX
if (financialAlert) decrease(stockOpts)
```

V2

Background and Motivation

- Variant 3: Multiple aggregators
 - **f1, f2, f3** are dispatched to n aggregators, that work autonomously
 - In case of deviating behaviour, any aggregator can adjust stockOpts
 - No glitch freedom required, but every single aggregator needs to see **f1, f2** and **f3** change in the same order


```
[var financialAlert_n := computeAlert_n(f1,f2,f3)  
| if (financialAlert_n) adjust_n(stockOpts)] v3
```

A model for DRP

- Formal definition of DRP system architecture/consistency guarantees
- **Components:** networked nodes in system

$$c_1 \dots c_n$$

- **Variables:** state of component c_i is represented by $V_i = \{v_{i1} : \tau_{i1} \dots v_{im} : \tau_{im}\}$



A model for DRP

- Besides traditional *imperative* variables, *reactive* and *observable* variables are defined
- **Reactive**: variable that is automatically updated based on reactive expression
- **Observable**: continuously changing var that is used to build expressions. Local or Global.
- e.g. stock market:
`f3 := Model3.compute(marketIndex, news)`


Reactive (+ observable) variable

Observable variables

A model for DRP

- **Dependency Graph:**

- Directed graph $D = \{V, E\}$, where V is the set of all observable/reactive variables and E is the set of all edges that connect directly depending variables
- E.g. stock market for Variant 1 + 2:



A model for DRP

- **Events:**
 - *Write* event: $w_x(v)$
 - Occurs when value x is written to variable v
 - *Read* event: $r_x(v)$
 - Occurs when value x is read from variable v
 - *Update* event: $u(S, w_x(v))$, $S = \{w_{y1}(v_1) \dots w_{yn}(v_n)\}$
 - Depending variable v is reactively update with value x due to the write events contained in the set S

A model for DRP

- **Consistency Guarantees**
 - **Exactly once delivery:** ensures that, in absence of failure, the communication channel does not lose or duplicate an update. More formally:

If $w_x(v)$ occurs, then $u(S_i, \bar{w}_y(v_i))$, $w_x(v) \in S_i$ occurs exactly once.

A model for DRP

- **Consistency Guarantees**
 - **FIFO ordering:** changes to a variable v in a component c are propagated to depending reactive expressions in the same order they occur in c . More formally:

$\forall v_i, v_j$, such that v_j depends on v_i , if $w_{x1}(v_i)$ occurs before $w_{x2}(v_i)$, then $u(S_1), w_{x1}(v_i) \in S_1$ occurs before $u(S_2), w_{x2}(v_i) \in S_2$

A model for DRP

- **Consistency Guarantees**

- **Causal ordering:** ensures that events that are causally connected occur in every component in the same order. More formally:

We define a *happened before* (\rightarrow) partial order relation:

- If two events e_1, e_2 , occur in the same process, then $e_1 \rightarrow e_2$ if and only if e_1 occurs before e_2
- If $e_1 = w_x(v_i)$ and $e_2 = u(S_i, w_y(v_j))$, $w_x(v_i) \in S_i$, then $e_1 \rightarrow e_2$ (a write happens before an update depending on it)
- If $e_1 \rightarrow e_2$ and $e_2 \rightarrow e_3$, then $e_1 \rightarrow e_3$ (transitivity)

- No guarantees are made for events that are not causally connected!

A model for DRP

- **Consistency Guarantees**
 - **Glitch freedom:** no partial updates due to propagation delays. More formally:

Consider the set V_d , containing all observable variables a reactive variable v depends on. Let us call $V_{d1} \subseteq V_d$ the set of variables that depend directly or indirectly from a variable v_1 . The update $u(S, w_x(v))$ is a *partial* update if $S \subset V_{d1}$. A glitch free system does not have partial updates.

A model for DRP

- **Consistency Guarantees**

- **Atomic consistency:** ensures that: (i) the system provides FIFO ordering, and (ii) every write event to an observable variable is atomically propagated to all (in)directly depending reactive variables.

More formally:

All the update events $u(S_i, w_y(v_i))$ triggered (directly or indirectly) by $w_x(v)$ are executed as a single operation

- This is stricter than glitch freedom

DREAM: API

- DREAM is entirely written in Java
- Observable variables → observable objects
 - Inherit from `Observable` abstract class
 - All non-void methods: *observable* methods
 - Generic method m that potentially changes return value of observable method obm : m impacts obm
 - Impacts should be known by runtime
 - Java Annotations

DREAM: API

- Example of observable class representing an integer:

```
1 public class ObservableInteger extends Observable {  
2     private int val;  
3  
4     // Constructors ...  
5  
6     @ImpactsOn(methods = { "get" })  
7     public final void set(int val) {  
8         this.val = val;  
9     }  
10  
11    public final int get() {  
12        return val;  
13    }  
14 }
```

DREAM: API

- Reactive variables → Reactive objects
- Created by using the `ReactiveFactory` class
 - Parses reactive expressions (strings with ANTLR)
 - Reactive objects can be observable (optional)
- Naming space:
 - Unique name: `c.obj.obm` for observable method `obm` of object `obj` in component `c`
 - For local objects: `obj.obm`

DREAM: API

- Example:


```
1 // Component c1
2 ObservableInteger obInt =
3   new ObservableInteger("obInt1", 1, LOCAL);
4 ObservableString obStr1 =
5   new ObservableString("obStr1", "a", GLOBAL);
6 ObservableString obStr2 = ...
7
8 // Component c2
9 ReactiveInteger rInt = ReactiveFactory.
10  getInteger("obInt.get() * 2");
11 ReactiveString rStr = ReactiveFactory.
12  getString("obStr1.get() + obStr2.get()");
13 while(true) {
14  System.out.println(rStr.get())
15  Thread.sleep(500)
16 }
17
18 // Component c3
19 ReactiveInteger strLen =
20  ReactiveFactory.getObservableInteger
21  ("c1.obString1.get().length()", "obString1Len");
```

DREAM: Implementation

- Architecture consists of two parts:
 - A client library on every component
 - A distributed event-based infrastructure,
consisting of *brokers*
- Brokers form an acyclic overlay network,
offering communication between components
- Optional registry for persistence

DREAM: Implementation

- Architecture overview



DREAM: Implementation

- **Pub-Sub Communication:**

Clients register with brokers through 3 primitives:

- `advertise(c,obj,obm)`: used by `c` if it has a globally observable method `obj.obm()`
- `subscribe(c,obj,obm)`: used to register a component that has a reactive expression containing `c.obj.obm()`
- `notify(c,obj,obm,val)`: used by `c` when `obj.obm()` has a new value `val`

DREAM: Implementation


- **Clients**
 - CommunicationManager:
 - Proxy for global communication
 - Manage local communication
 - Observable objects:
 - Have Watcher code woven in through AOP
 - Watcher interacts with CommunicationManager to:
 1. Advertise new objects through `advertise(c,obj,obm)`
 2. Detect changes to observables and propagate them out through `notify(c,obj,obm,val)`

DREAM: Implementation

- **Clients**
 - Reactive Objects:
 - When instantiated, for all relevant observable methods
→ `subscribe(c,obj,obm)` with `CommunicationManager`
 - When new values available, notification from `CommunicationManager`

DREAM: Implementation

- Architecture overview



DREAMS: Implementation

- **Brokers**

- Run REDS event dispatching

- Brokers are connected in acyclic graph
 - Advertisements are propagated through graph + stored by all brokers, remembering next hop
 - When a broker receives a subscription, store in table and forward to next hop (retrace path of advertisements)

DREAMS: Implementation

- **Consistency Guarantees**
 - Causal ordering:
 - Use point to point TCP for broker-broker and client-broker communication
 - Use single thread for FIFO event processing
- These 2 properties with an acyclic topology are sufficient for causal ordering

DREAMS: Implementation

- **Consistency Guarantees**

- Glitch freedom:

- New reactive object: push propagate expression to *all* brokers → each broker has dependency graph
 - When a chain of operations is triggered, always include the original write event that caused it in communications
 - Local communication *has* to go through a broker as well to ensure glitch freedom
 - This information is enough for the brokers to schedule propagation in a way that avoids partial updates

DREAMS: Implementation

- **Consistency Guarantees**

- Atomic ordering:

- Adds centralized Ticket Granting Service (TGS)
 - When a write event occurs, *all* its directly and indirectly dependent reactive expressions are reevaluated atomically (no other write operations)
 - On write: get ticket, wait in line and be served one at a time
- This entails glitch freedom and is an even stronger consistency guarantee

Evaluation

- Twofold:
 1. Large scale emulation: Cost of DRP protocols with different levels of consistency guarantees/ varying parameters. KPIs:
 - Average propagation delay (ms)
 - Network wide traffic throughput (KB/s)
 2. Real-world runtime overheads

Evaluation

- Default values for emulation:

Number of brokers	10
Number of components	50
Topology of broker network	Scale-free
Percentage of pure forwarders	50%
Distribution of components	Uniform
Link latency	1 ms–5 ms
Number of reactive graphs	10
Size of dependency graphs	5
Size of reactive expressions	2
Degree of locality in expressions	0.8
Frequency of change for observable objects	1 change/s


Evaluation

- **Advantages of distribution**
 - 1 broker vs. 10 brokers
 - Causal: no big impact – mainly due to locality
 - Glitch free: *all* propagation through broker
 - Having multiple brokers helps
 - Atomic: adds TGS delay + traffic
 - Same advantages when multiple brokers

	Delay (ms)		Traffic (KB/s)	
	Centr.	Distr.	Centr.	Distr.
Causal	4.77	4.76	68.3	69.8
Glitch free	29.53	17.18	205.4	130.9
Atomic	53.41	26.75	265.5	161.3


Evaluation

- **Locality of expressions**
 - General trend: locality cuts costs



(a) Delay

Completely remote + causal = glitches!




(b) Traffic

Completely local + causal = 0 costs


Evaluation

- **Size of reactive graphs**
 - General trend: large reactive graphs increase costs

Long chains of reactive vars + causal = glitches!



(a) Delay




(b) Traffic


Evaluation

- **Size of expressions**
 - General trend: bigger expressions increase costs

More vars/expression + causal = glitches!



(a) Delay




(b) Traffic

Evaluation

- **Runtime overheads**
 - Overheads consisting of:
 - Intercepting a method call
 - Serializing/deserializing
 - Propagating the change
 - Evaluating reactive expression
 - Local scenario: two clients and a broker on 1 machine, with increasing expression length

Evaluation

- **Runtime overheads**
 - Conclusion: runtime overheads are minimal



Conclusion

- Key contributions:
 - First abstract model of DRP/formalizing consistency constraints
 - DREAM: a first DRP middleware supporting 3 propagation semantics
 - A thorough evaluation of the costs

Conclusion

- Future work:
 - A glitch free protocol that takes advantage of locality
 - Robustness in case of node failure
 - More complex expressions (time series and sequence of changes)
 - Different evaluation strategies (lazy, incremental) to improve efficiency
 - More real applications